

Model-Parallel Deep Learning

Efficient DL, Episode ++i '26

Yandex
Research



Dealing with large models

~~Model-Parallel Deep Learning~~

Efficient DL, Episode ++i '26

Yandex
Research



Recap: large models

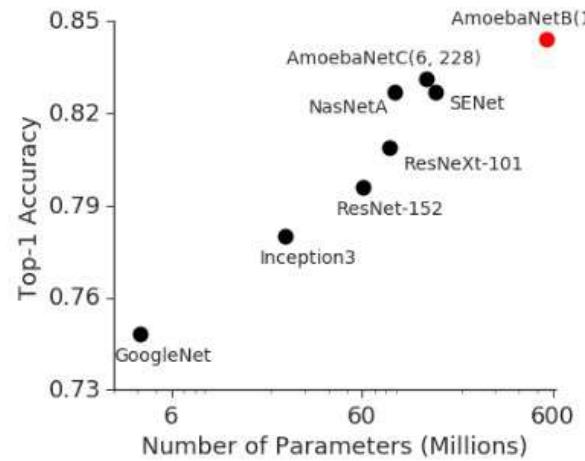
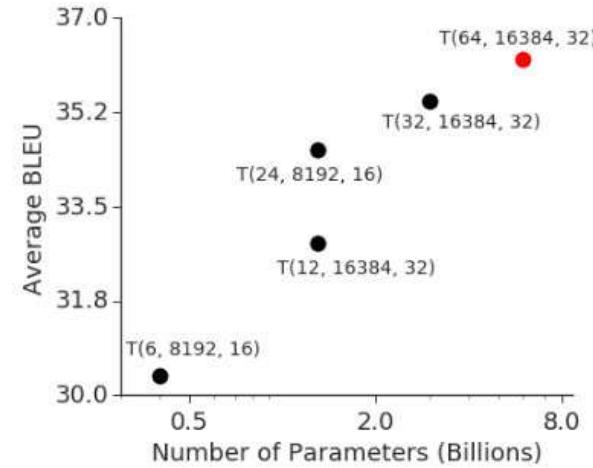


Image Classification
ImageNet

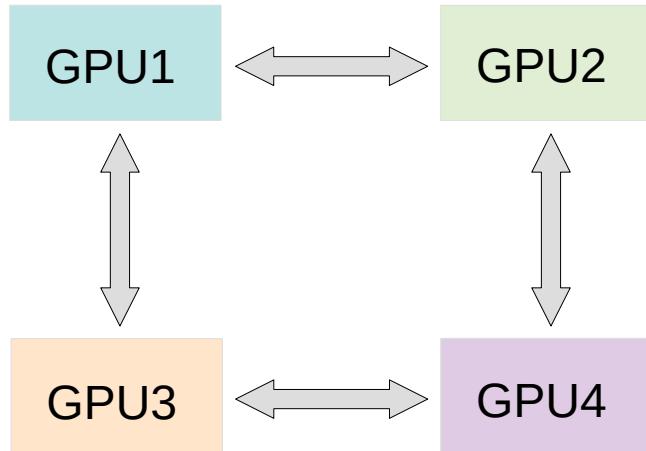


Machine Translation
average over WMT

Source: <https://arxiv.org/abs/1811.06965>

Recap: Ring allreduce

Bonus quest: you can only send data between **adjacent** gpus



Ring topology



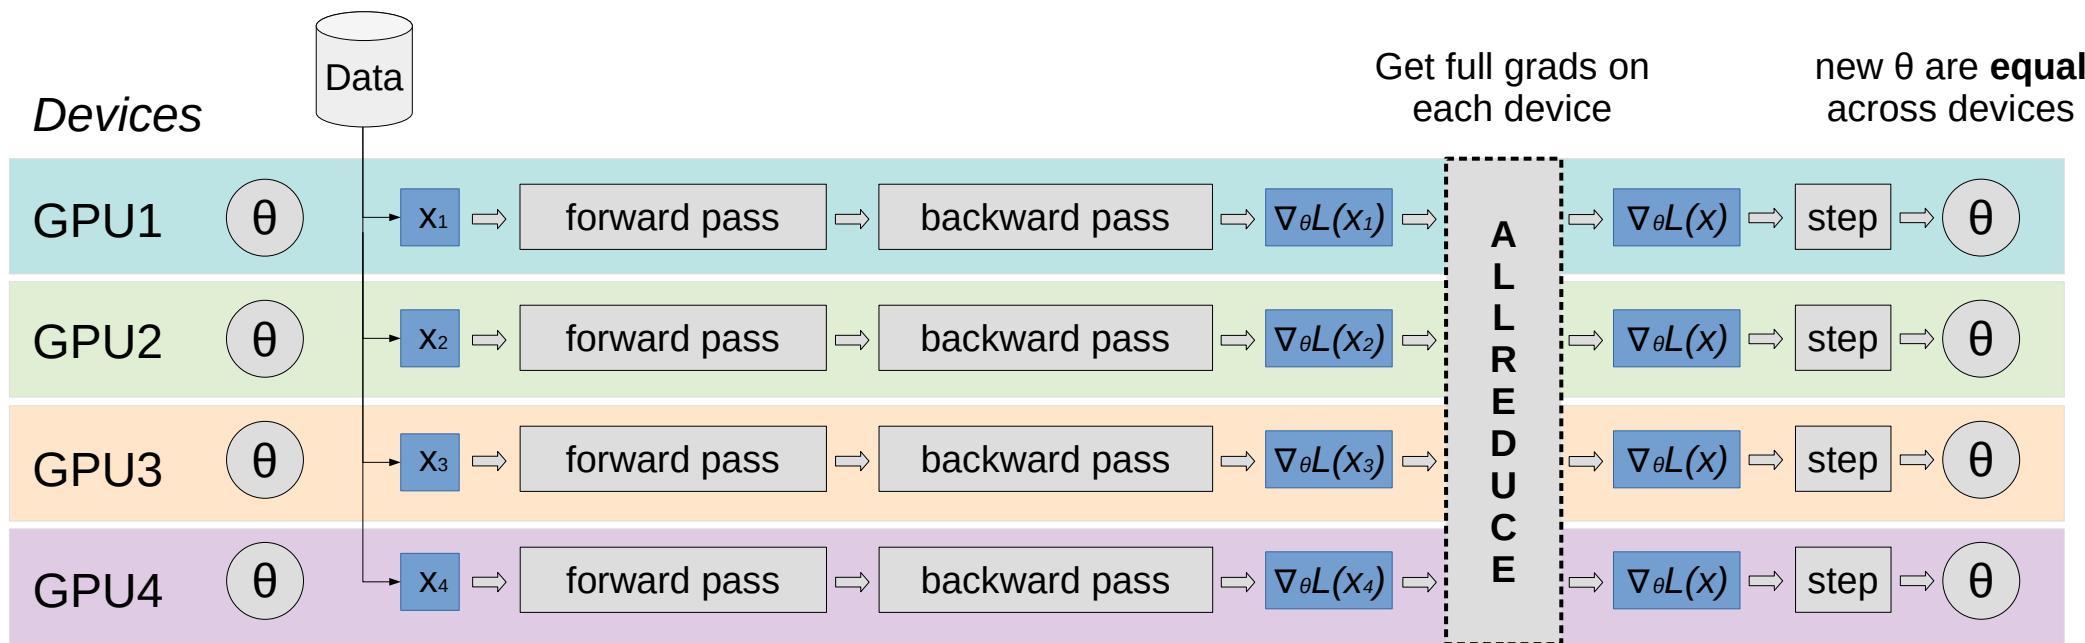
Image: graphcore ipu server

Answer & more: tinyurl.com/ring-allreduce-blog

Recap: All-Reduce SGD

arxiv.org/abs/1706.02677

Idea: get rid of the host, each gpu runs its own computation
Q: why will weights be equal after such step?



Q: What if a model is larger than GPU?

Q: What if a model is larger than GPU?
easy mode: cannot fit the right batch size

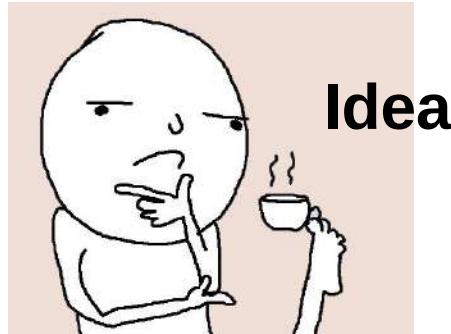
hard mode: cannot fit a single sample

expert mode: not even parameters!

Q: What if a model is larger than GPU?
easy mode: cannot fit the right batch size

hard mode: cannot fit a single sample

expert mode: not even parameters!



Ideas?

Q: What if a model is larger than GPU?

easy mode: cannot fit the right batch size

hard mode: cannot fit a single sample

expert mode: not even parameters!

Solution: accumulate
grads from several
training batches

```
[ ] 1 optimizer.zero_grad()  
2 for i in range(B):  
3     loss = model(**next_batch())  
4     (loss / B).backward()  
5 optimizer.step()
```

Q: What if a model is larger than GPU?

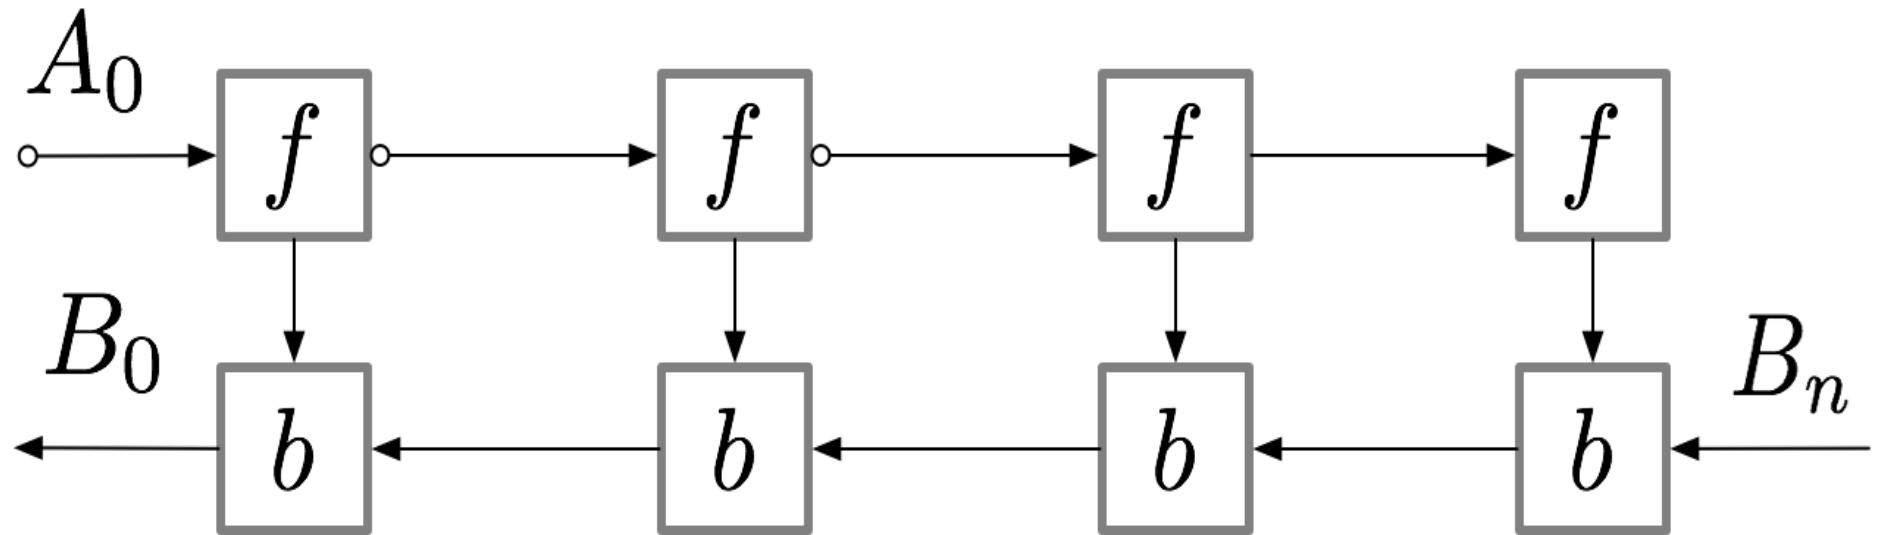
easy mode: cannot fit the right batch size

hard mode: cannot fit one training sample

expert mode: not even parameters!

Gradient checkpointing

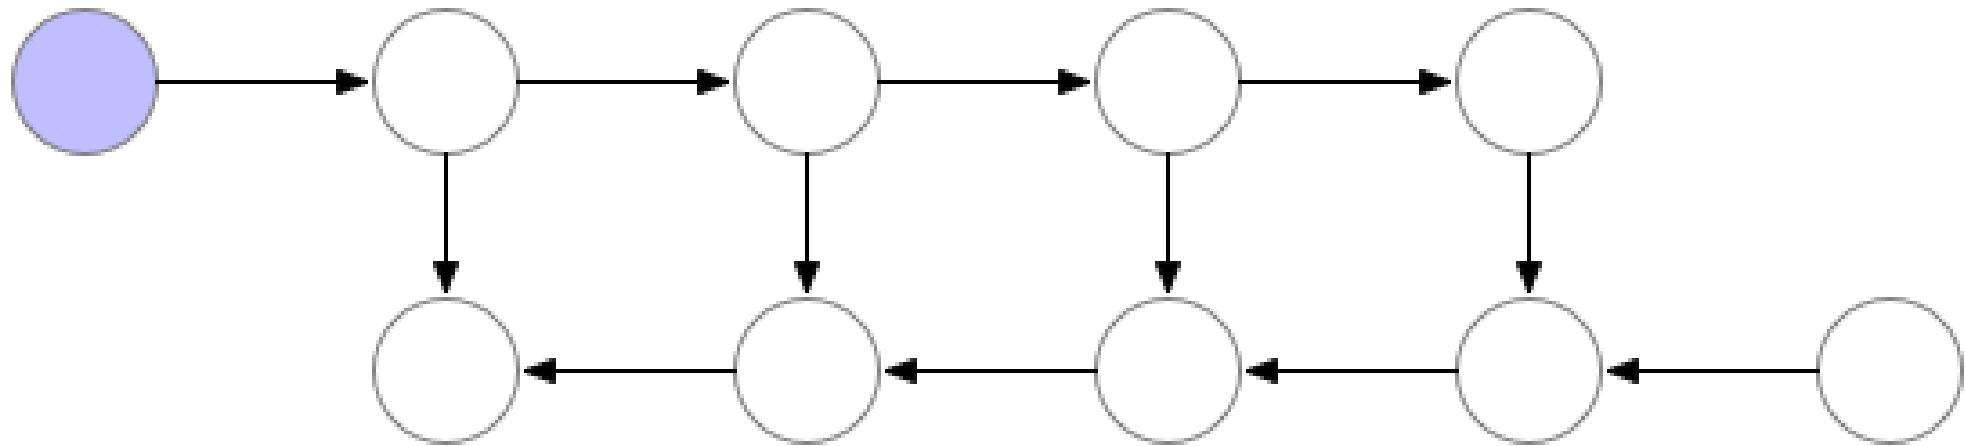
aka rematerialization



Paper (DL): arxiv.org/pdf/1604.06174.pdf
TF: github.com/cybertronai/gradient-checkpointing
Pytorch: pytorch.org/docs/stable/checkpoint.html

Gradient checkpointing

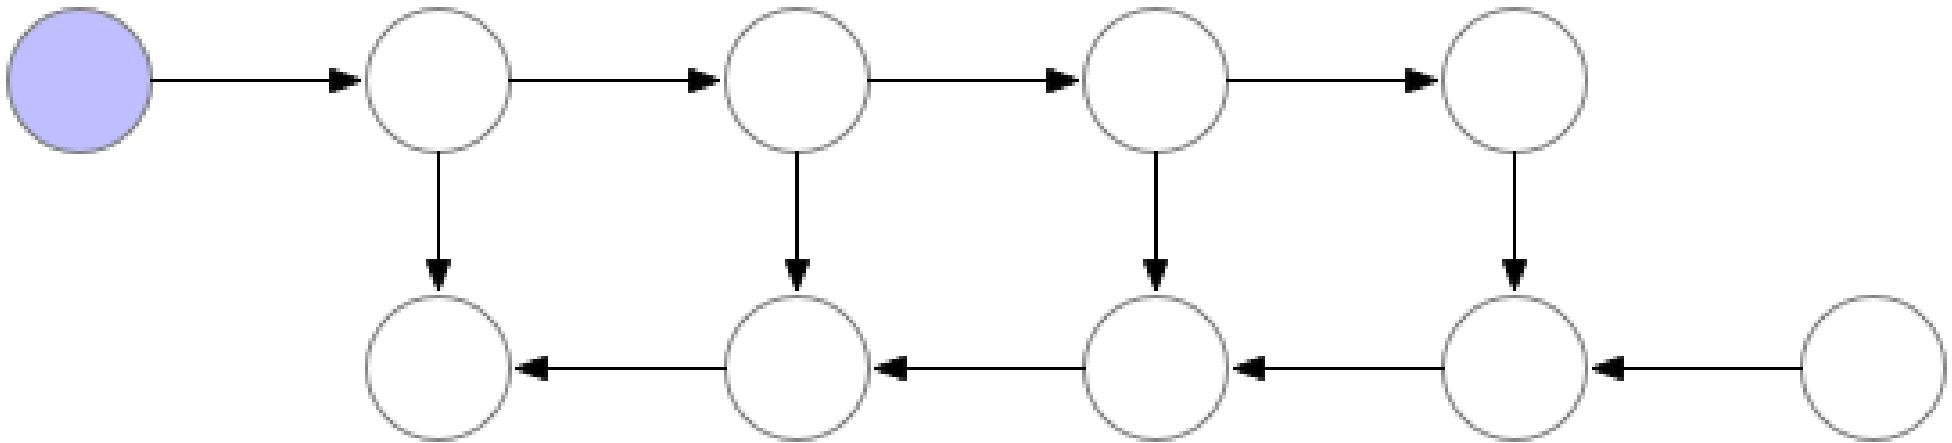
Normal backprop



Paper (DL): arxiv.org/pdf/1604.06174.pdf
TF: github.com/cybertronai/gradient-checkpointing
Pytorch: pytorch.org/docs/stable/checkpoint.html

Gradient checkpointing

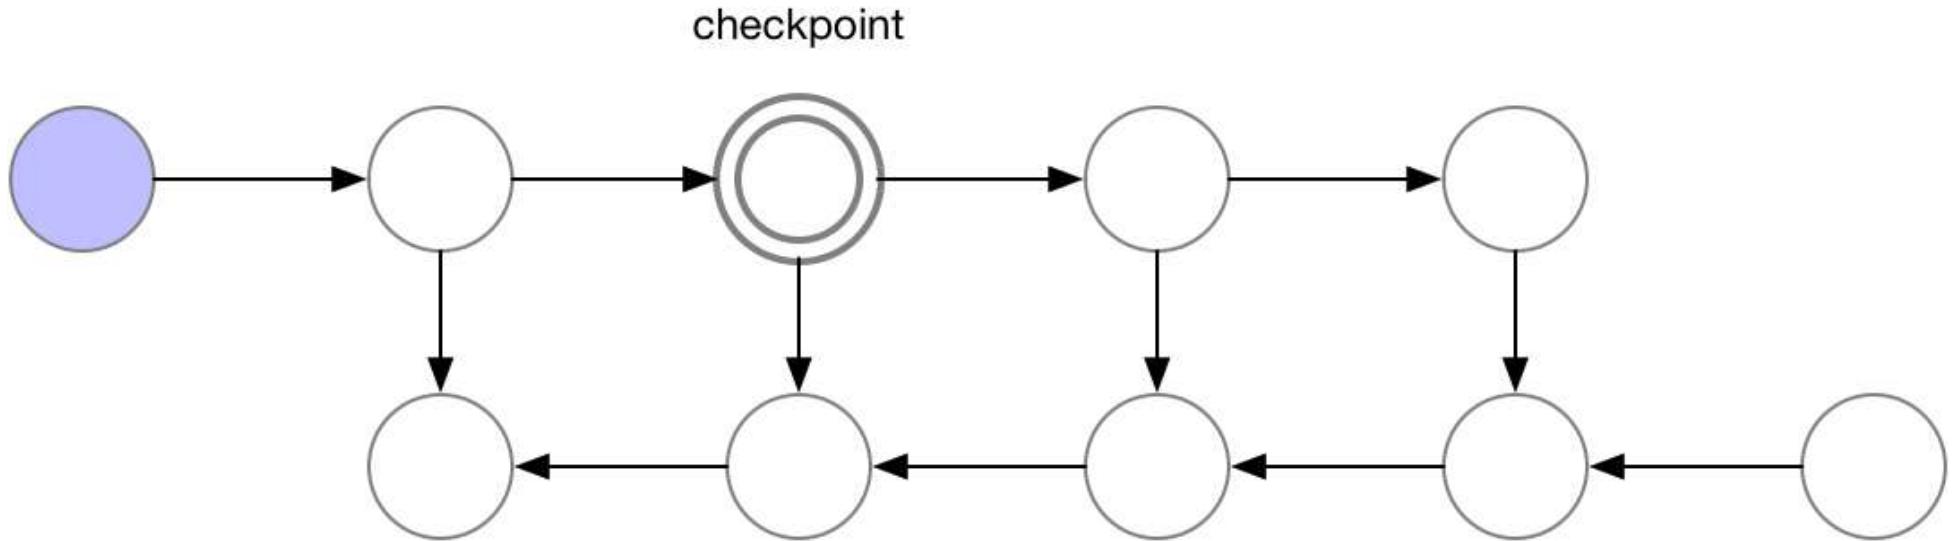
Full rematerialization



Paper (DL): arxiv.org/pdf/1604.06174.pdf
TF: github.com/cybertronai/gradient-checkpointing
Pytorch: pytorch.org/docs/stable/checkpoint.html

Gradient checkpointing

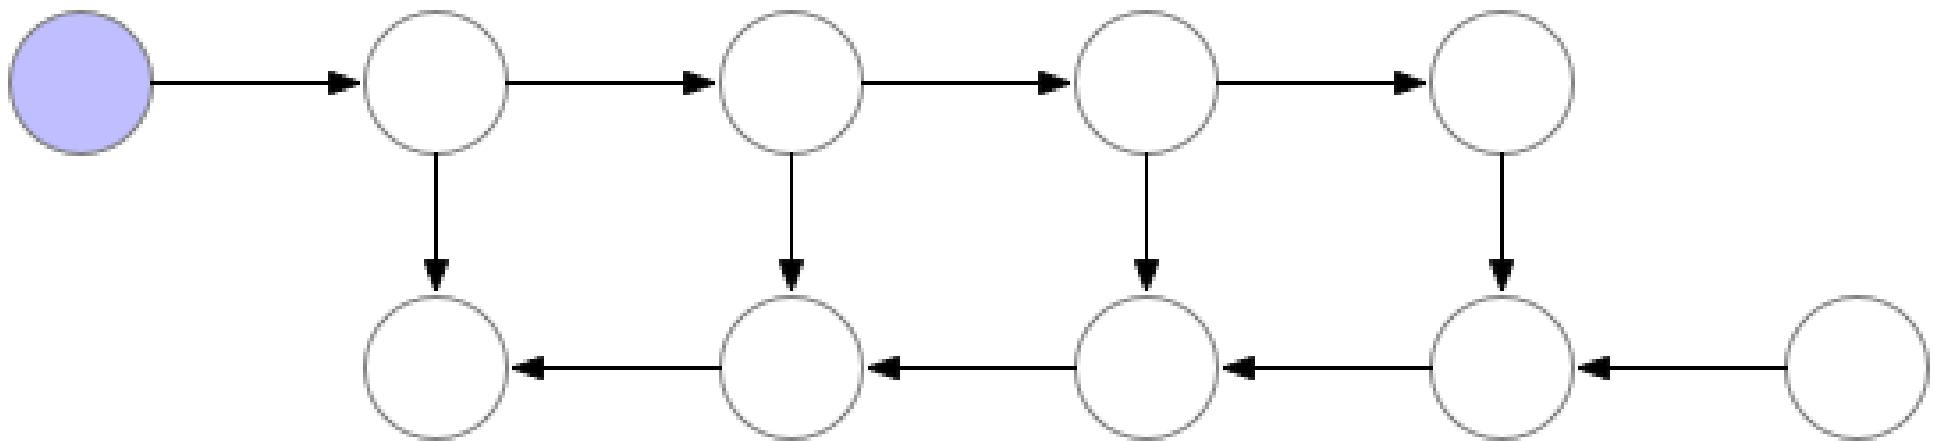
Single checkpoint



Paper (DL): arxiv.org/pdf/1604.06174.pdf
TF: github.com/cybertronai/gradient-checkpointing
Pytorch: pytorch.org/docs/stable/checkpoint.html

Gradient checkpointing

Single checkpoint



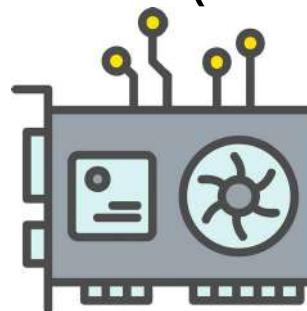
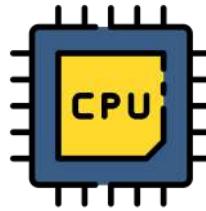
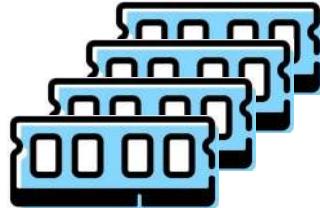
Paper (DL): arxiv.org/pdf/1604.06174.pdf
TF: github.com/cybertronai/gradient-checkpointing
Pytorch: pytorch.org/docs/stable/checkpoint.html

Q: What if a model is larger than GPU?

easy mode: cannot fit batch size 1

expert mode: not even parameters!

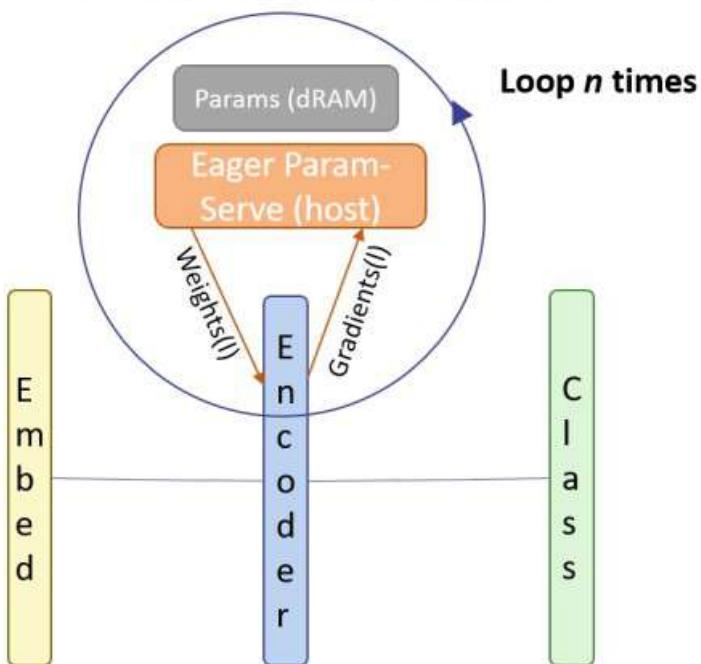
You still have one GPU... (*but not only a GPU*)



Memory offloading

L2L: <https://arxiv.org/abs/2002.05645>

EPS with L2L execution



- Initialize all layers on CPU
- Move k layers at a time to GPU
- Remove layers after computation
- Fetch k+1-st layer while k-th runs
- Still 20-50% overhead

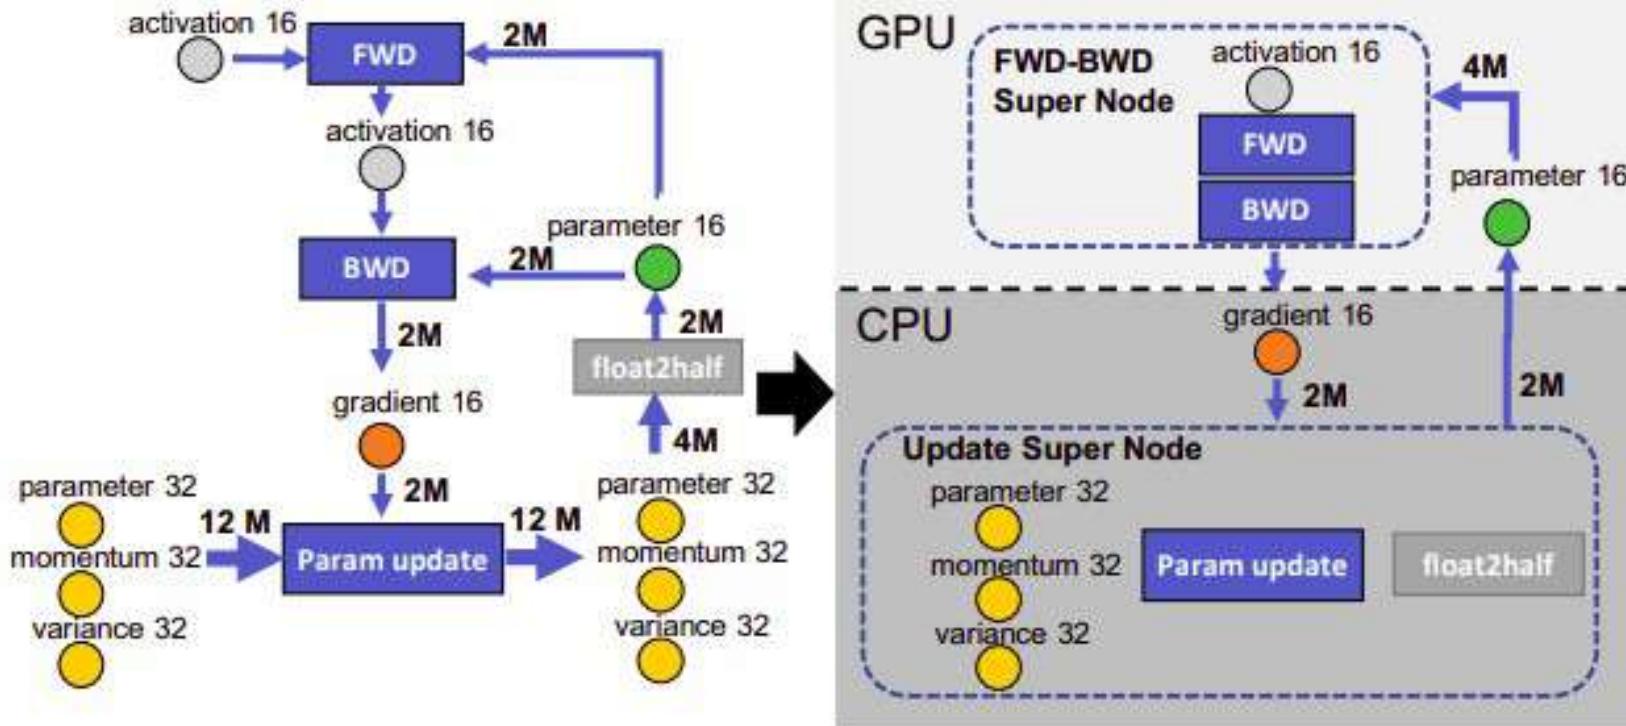
Memory offloading

L2L: <https://arxiv.org/abs/2002.05645>

METHOD	UBATCH SIZE	DEVICE BATCH SIZE	#LAYER	#PARAMETERS	MEMORY (GB)
BASELINE	2	2	24	300 MILLION	9.23
BASELINE	2	2	48	600 MILLION	OOM
L2L-STASH ON GPU	64	64	24	300 MILLION	5.22
L2L-STASH ON GPU	64	64	48	600 MILLION	6.76
L2L-STASH ON GPU	64	64	96	1.2 BILLION	9.83
L2L-STASH ON CPU	64	64	24	300 MILLION	3.69
L2L-STASH ON CPU	64	64	96	1.2 BILLION	3.69
L2L-STASH ON CPU	64	64	384	4.8 BILLION	3.69

Memory offloading

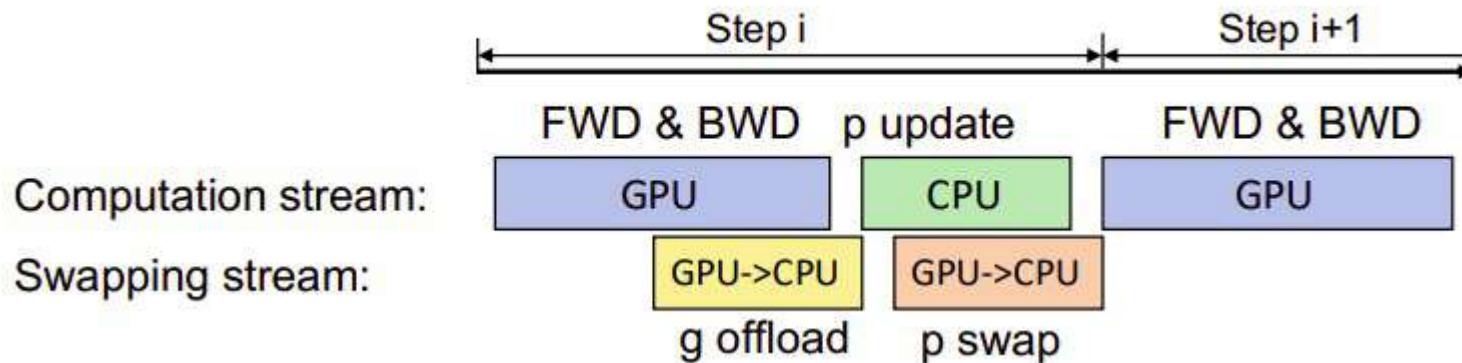
ZeRO-offload: <https://arxiv.org/abs/2101.06840>



Memory offloading

ZeRO-offload: <https://arxiv.org/abs/2101.06840>

- Offload **in parallel** with computation
- Use gradient checkpointing
- Delayed parameter update



Memory offloading

ZeRO-offload: <https://arxiv.org/abs/2101.06840>

- Offload in parallel with computation
- Use gradient checkpointing
- **Delayed parameter update**

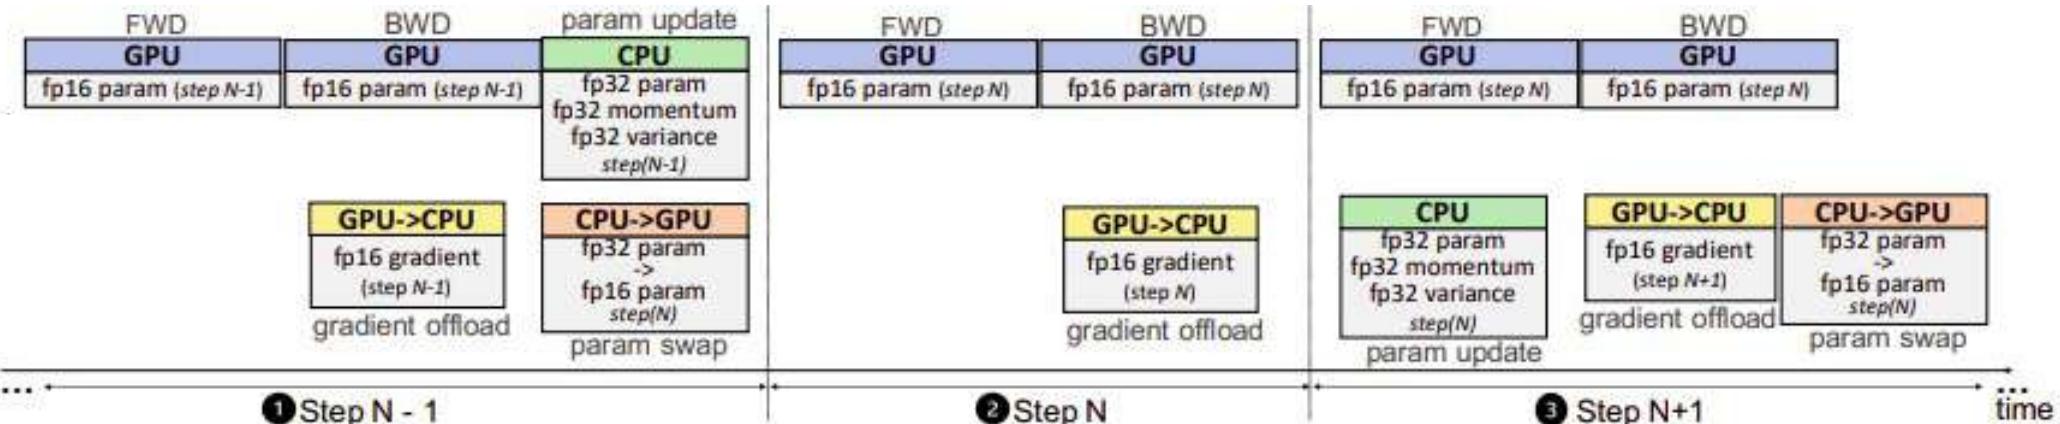
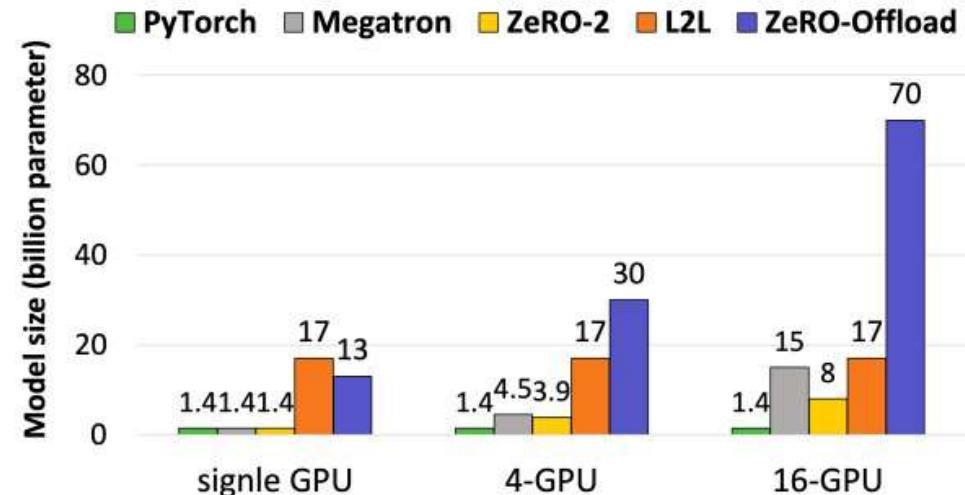
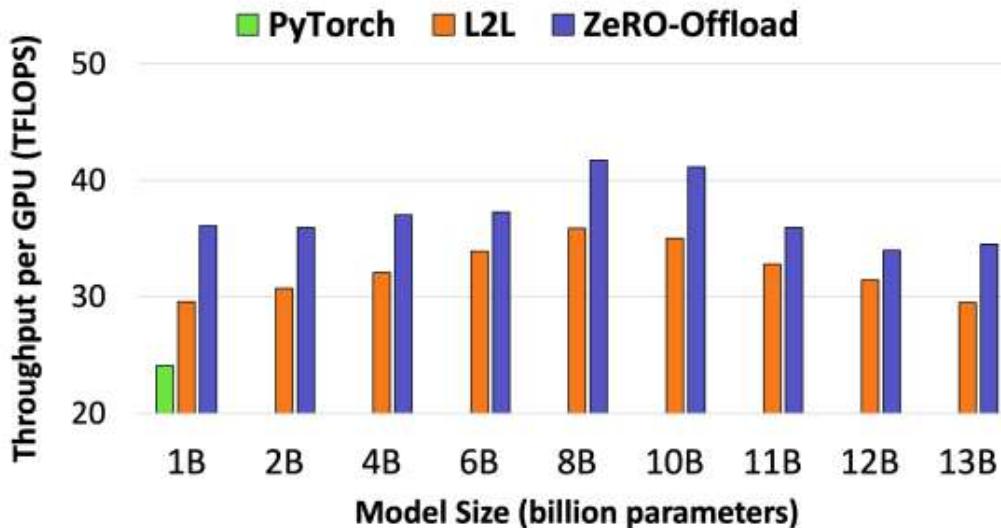


Figure 6: Delayed parameter update during the training process.

Memory offloading

ZeRO-offload: <https://arxiv.org/abs/2101.06840>

- Offload in parallel with computation
- Use gradient checkpointing
- Delayed parameter update

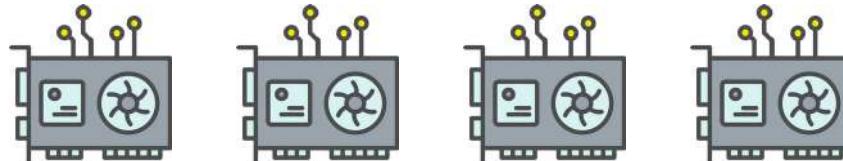


Q: What if a model is larger than GPU?

easy mode: cannot fit batch size 1

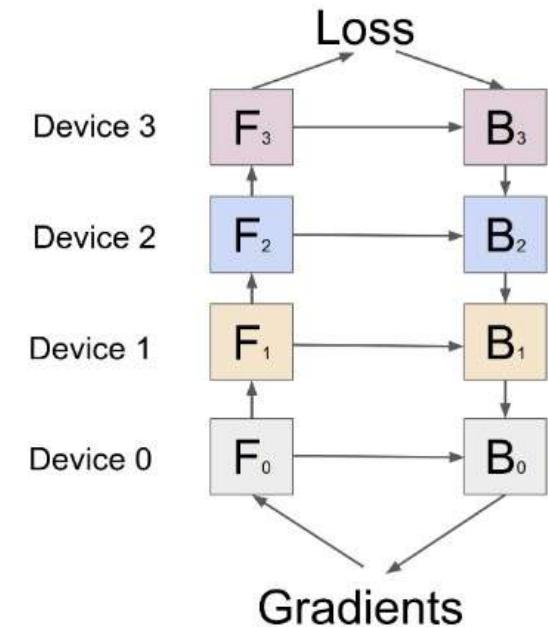
expert mode: not even parameters!

**Can we do it better with
multiple GPUs?**



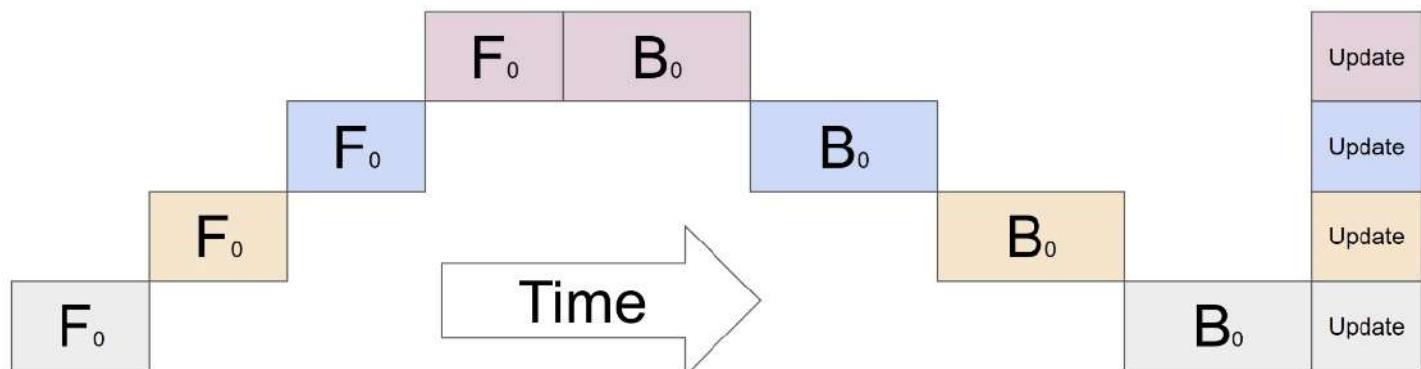
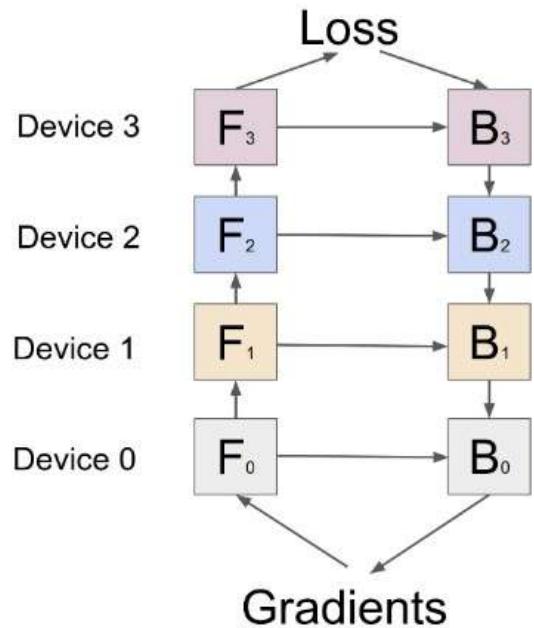
Model-parallel training

Q: What if a model is larger than GPU?



Model-parallel training

Q: What if a model is larger than GPU?



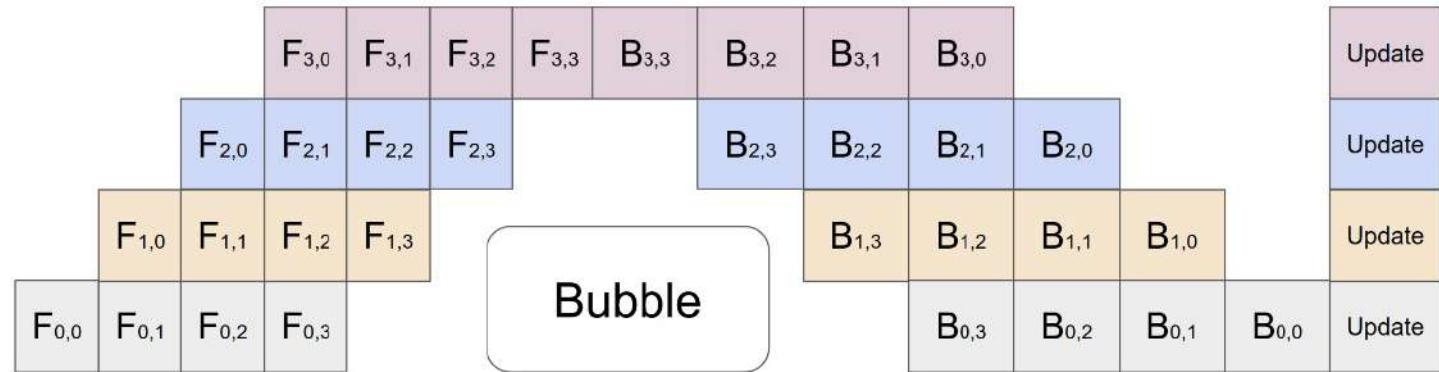
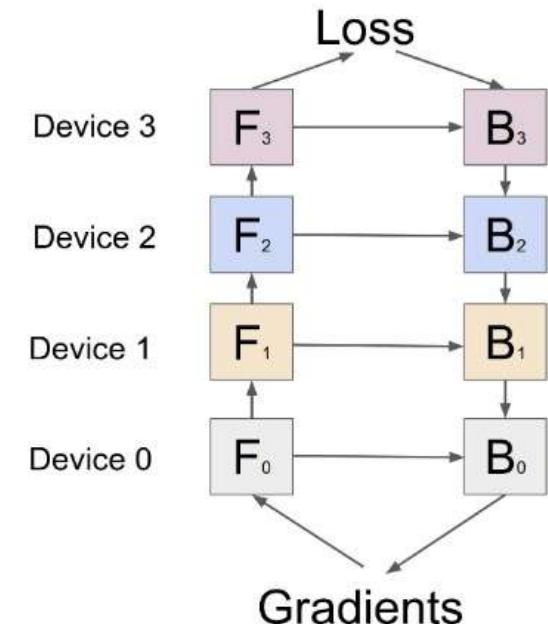
model size: O(N)
throughput: O(1)

Q: Can we go faster?

Pipelining

GPipe: arxiv.org/abs/1811.06965 – good starting point, *not* the 1st paper

Idea: split data into micro-batches and form a pipeline (right)

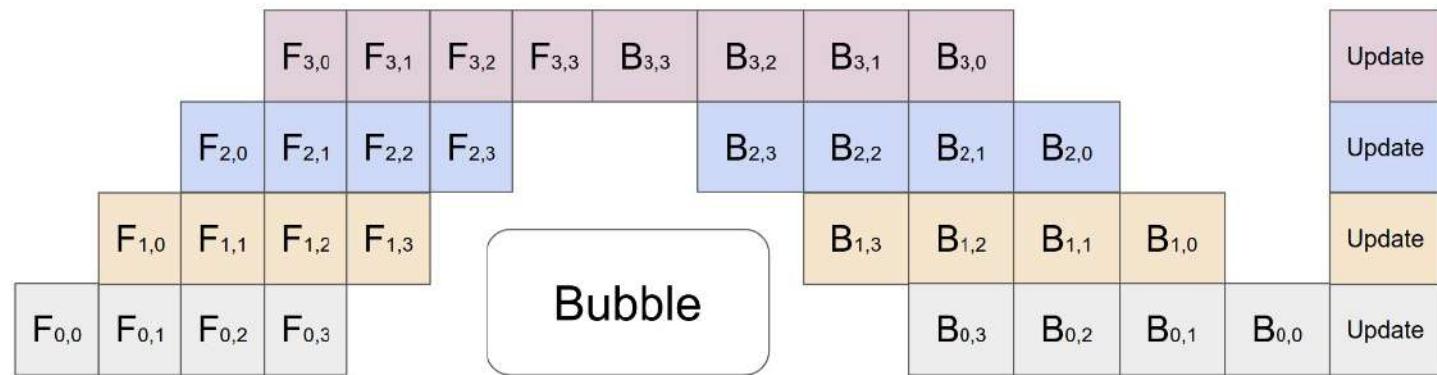
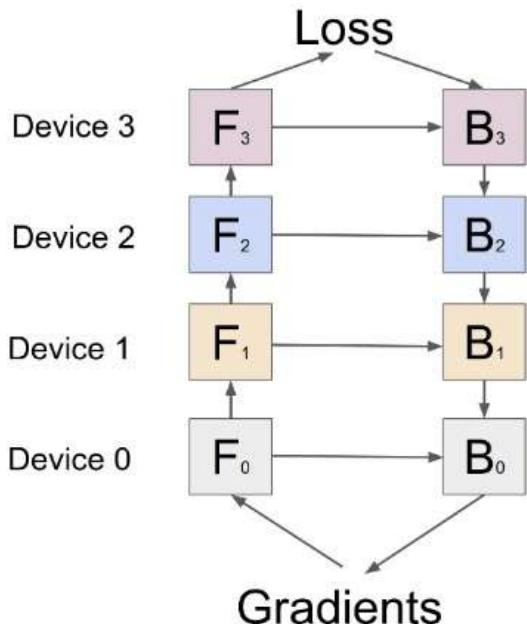


model size: $O(n)$
throughput: $O(n)$ – with caveats

Pipelining

GPipe: arxiv.org/abs/1811.06965 – good starting point, *not* the 1st paper

Idea: split data into micro-batches and form a pipeline (right)



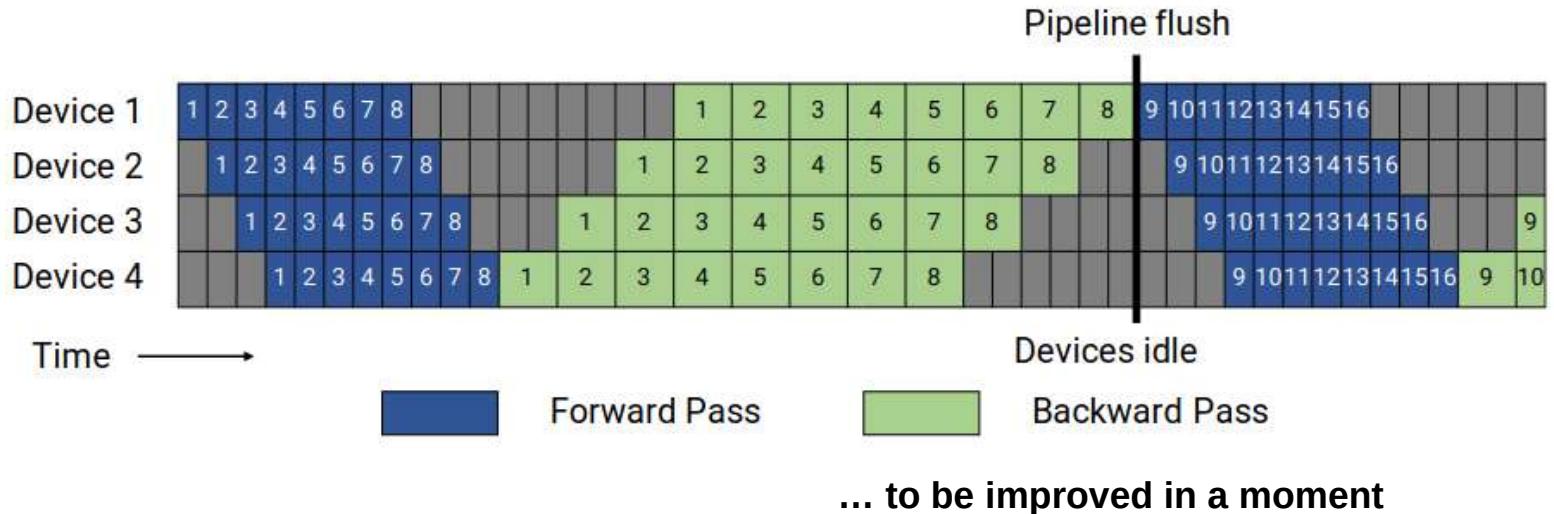
model size: $O(n)$
throughput: $O(n)$ – with caveats

Q: Even faster?

Reducing the bubble

GPipe: arxiv.org/abs/1811.06965

GPipe:



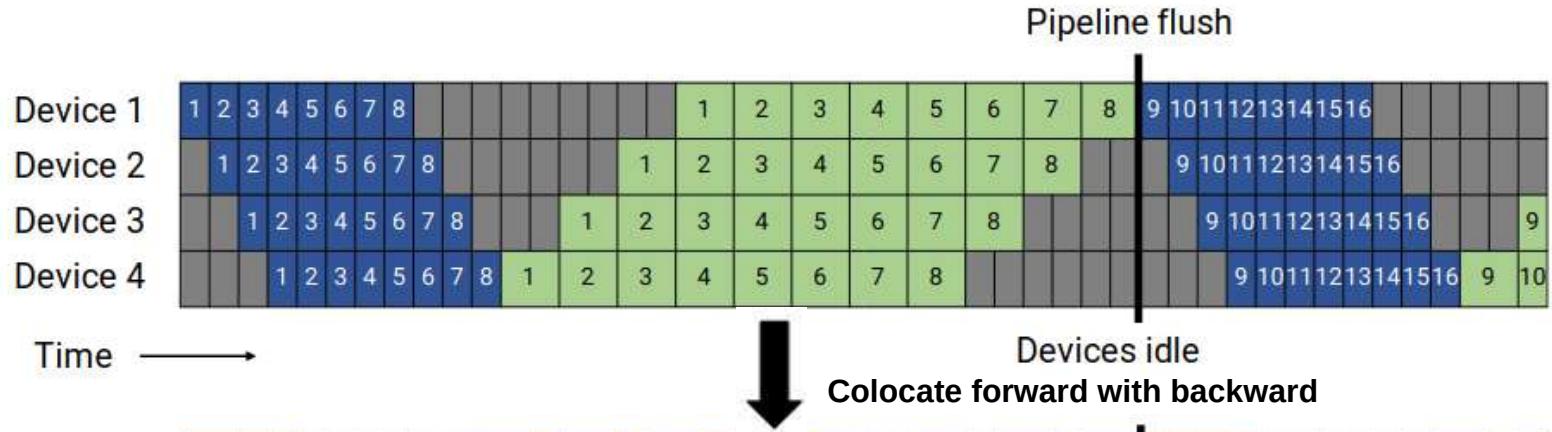
Note: backward takes longer than forward in practice

E.g. linear forward has one matmul,
backward has two matmuls (dW and dX)

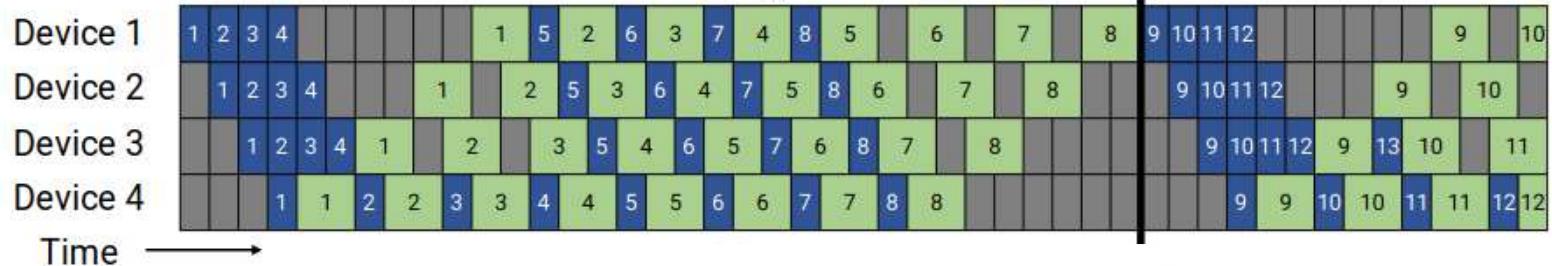
Reducing the bubble

1F1B pipeline from Megatron: <https://arxiv.org/abs/2104.04473>

GPipe:



1F1B:



Forward Pass

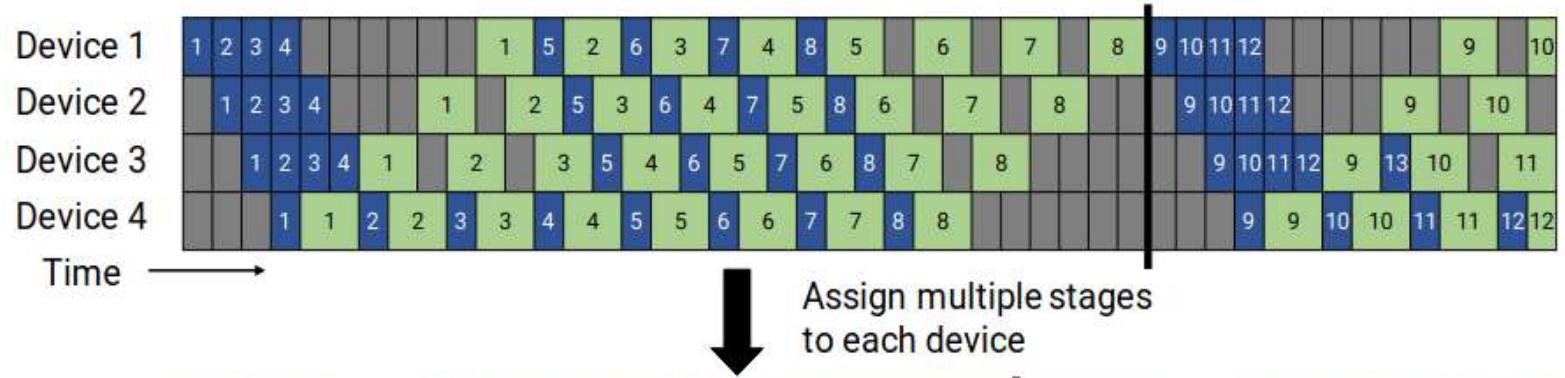


Backward Pass

Reducing the bubble (further)

1F1B pipeline from Megatron: <https://arxiv.org/abs/2104.04473>

1F1B:

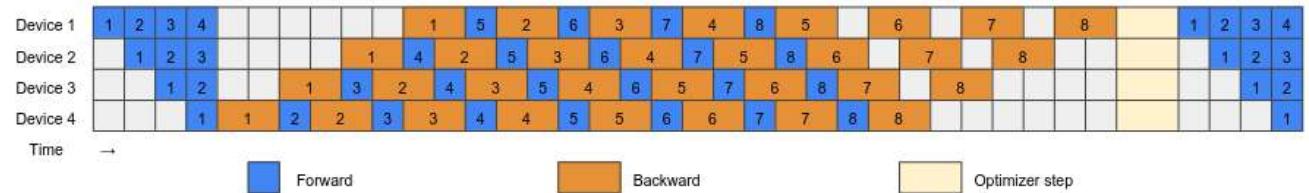


1F1B interleaved:



Reducing the bubble (furtherer)

ZB1P: “almost zero bubble” <https://arxiv.org/abs/2401.10241>



Idea: split backward into two ops:
- w.r.t inputs and w.r.t. weights

Grad w.r.t. weights doesn't block
backward pass to prev stage

Figure 2: 1F1B pipeline schedule.

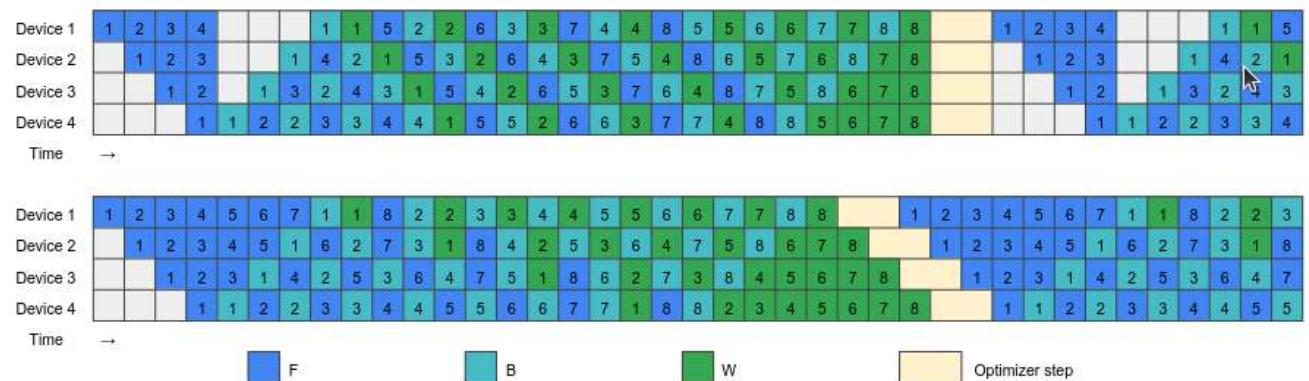


Figure 3: Handcrafted pipeline schedules, top: ZB-H1; bottom: ZB-H2

Reducing the bubble (furtherer yet)

Deepseek V1 schedule: <https://arxiv.org/abs/2412.19437>

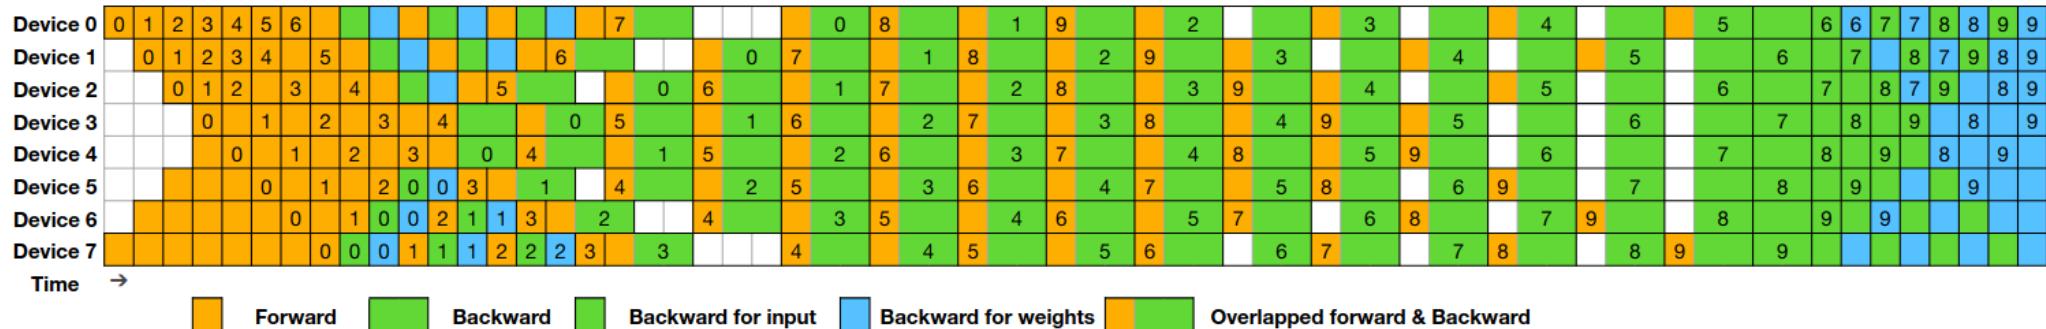


Figure 5 | Example DualPipe scheduling for 8 PP ranks and 20 micro-batches in two directions. The micro-batches in the reverse direction are symmetric to those in the forward direction, so we omit their batch ID for illustration simplicity. Two cells enclosed by a shared black border have mutually overlapped computation and communication.

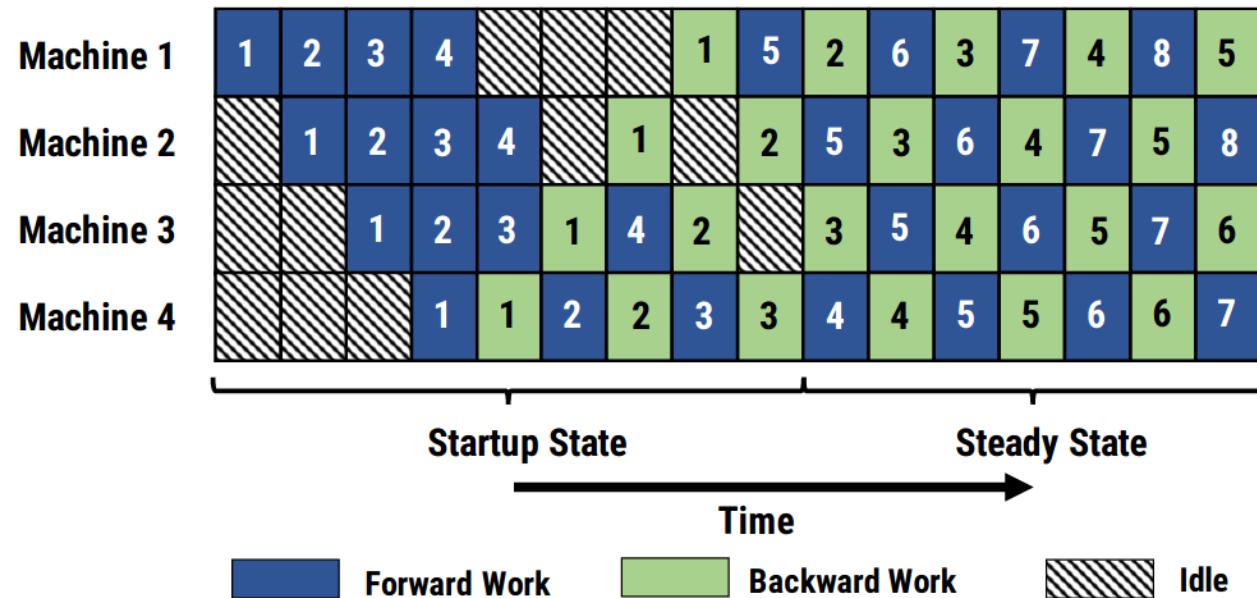
Asynchronous Pipelining

PipeDream: arxiv.org/abs/1806.03377

Idea: apply gradients with every microbatch for maximum throughput

Also neat:

- Automatically partition layers to GPUs via dynamic programming
- Store k past weight versions to reduce gradient staleness
- Aims at high latency



Pipelining Recap

When to use:

- model doesn't fit on GPU; have multiple GPUs
- if model fits, but not the activations: ???
- if model doesn't fit, but you only have one GPU: ???

How to use:

(just a moment...)

Pipelining Recap

When to use:

- model doesn't fit on GPU; have multiple GPUs
- if model fits, but not the activations: just do grad checkpointing!
- if model doesn't fit, but you only have one GPU: offloading!

How to use:

- Basic implementation (GPipe): github.com/kakaobrain/torchgpipe

```
from torchgpipe import GPipe
model = nn.Sequential(a, b, c, d)
model = GPipe(model, balance=[1, 1, 1, 1], chunks=8)
output = model(input)
```



Pipelining Recap

When to use:

- model doesn't fit on GPU; have multiple GPUs
- if model fits, but not the activations: just do grad checkpointing!
- if model doesn't fit, but you only have one GPU: offloading!

How to use:

- Basic implementation (GPipe): github.com/kakaobrain/torchgpipe
- PyTorch built-in: pytorch.org/tutorials/intermediate/pipelining_tutorial.html

```
from torch.distributed.pipeline import ScheduleGPipe  
  
# Create a schedule  
schedule = ScheduleGPipe(stage, n_microbatches)
```



uses `torch.distributed` (`torchrun`) | supports GPipe, 1F1B, extendable!

Pipelining Recap

When to use:

- model doesn't fit on GPU; have multiple GPUs
- if model fits, but not the activations: just do grad checkpointing!
- if model doesn't fit, but you only have one GPU: offloading!

How to use:

- Basic implementation (GPipe): github.com/kakaobrain/torchgpipe
- PyTorch built-in: pytorch.org/tutorials/intermediate/pipelining_tutorial.html
- DeepSpeed: <https://deepspeed.readthedocs.io/en/latest/pipeline.html>

Custom pipelines in many applications

- Megatron-LM: <https://github.com/NVIDIA/Megatron-LM> (transformer-specific)
- Megablocks: <https://github.com/databricks/megablocks> (mixture-of-experts)

Pipelining Recap

When to use:

- model doesn't fit on GPU; have multiple GPUs
- if model fits, but not the activations: just do grad checkpointing!
- if model doesn't fit, but you only have one GPU: offloading!

How to use:

- Basic implementation (GPipe): github.com/kakaobrain/torchgpipe
- PyTorch built-in: pytorch.org/tutorials/intermediate/pipelining_tutorial.html
- DeepSpeed: <https://deepspeed.readthedocs.io/en/latest/pipeline.html>

Problems:

- Bubbles = wasted compute time (duh)
- What if model layers aren't symmetric? (e.g. LLM "head", local attn, ViT pooling)
Balancing a pipeline is a world of hurt.

[short break]
**How else can we run a large model
over multiple GPUs / hosts?**

Tensor-parallel training

<https://papers.nips.cc/paper/4824-imagenet-classification-with-deep-convolutional-neural-networks>

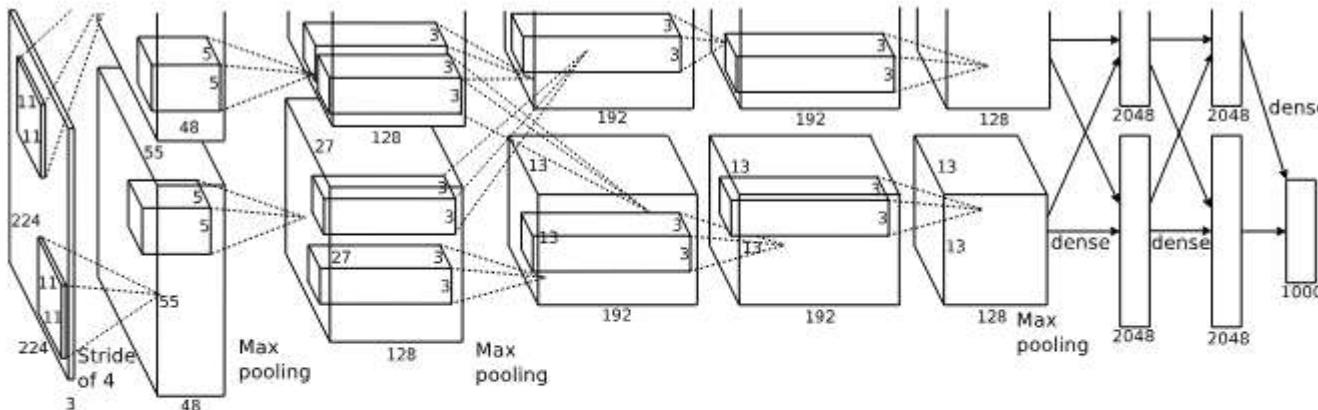
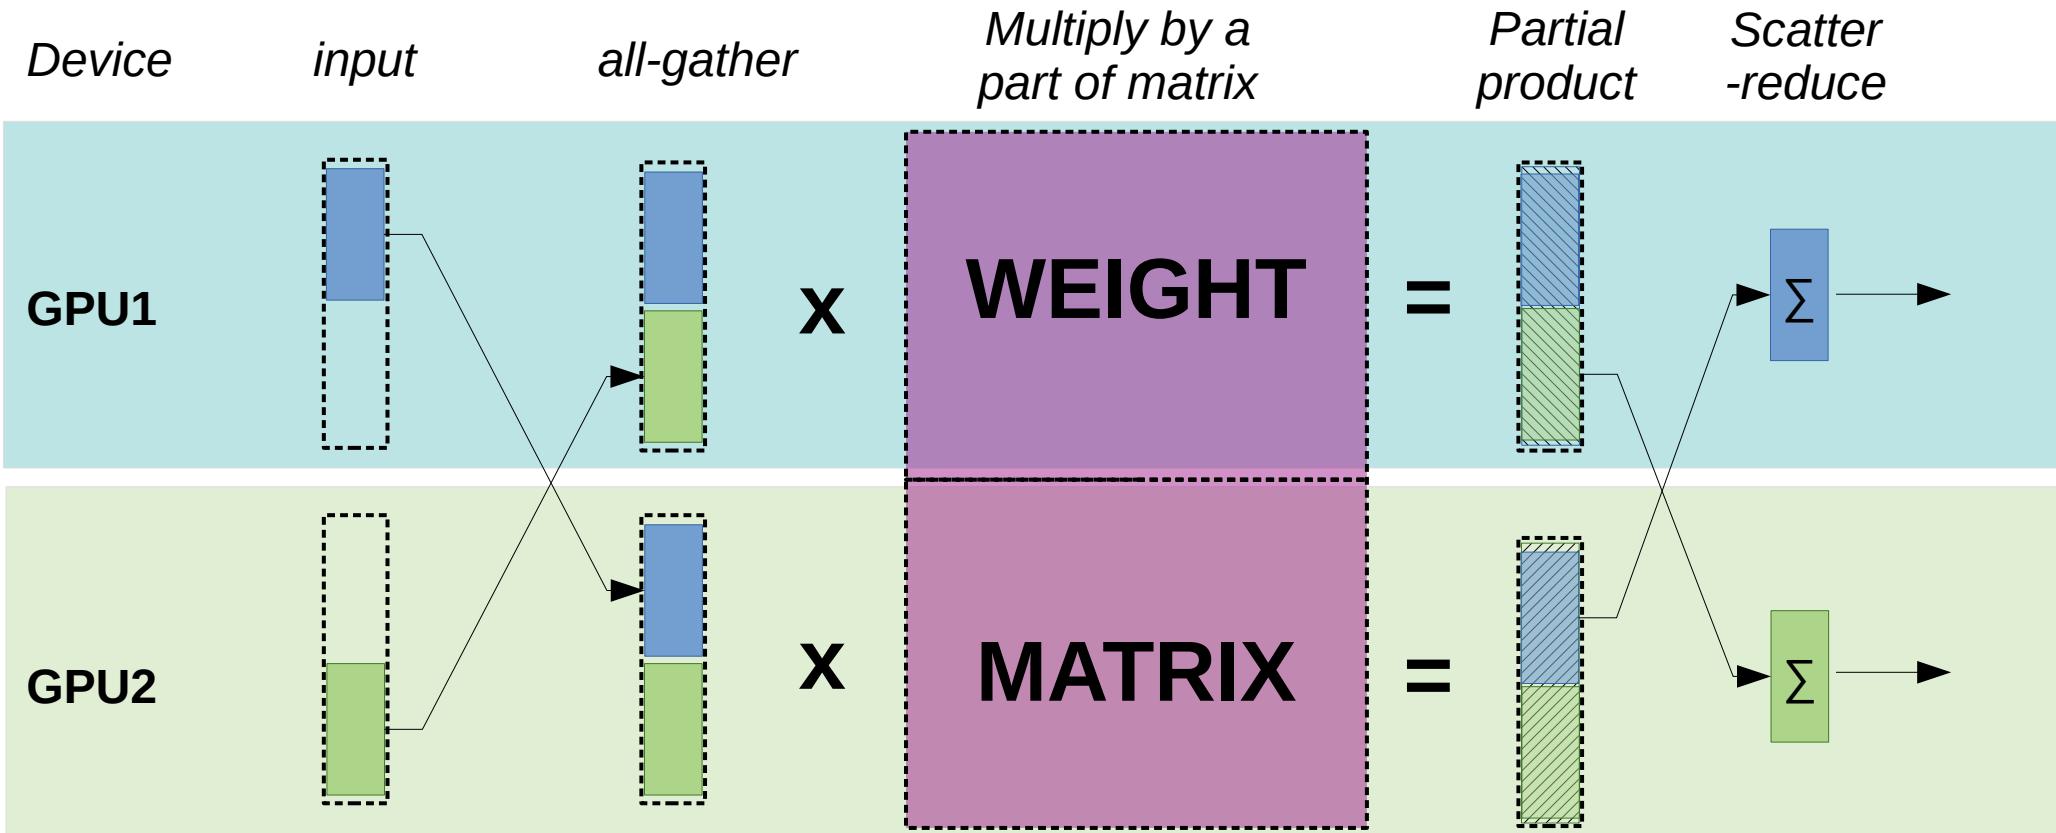


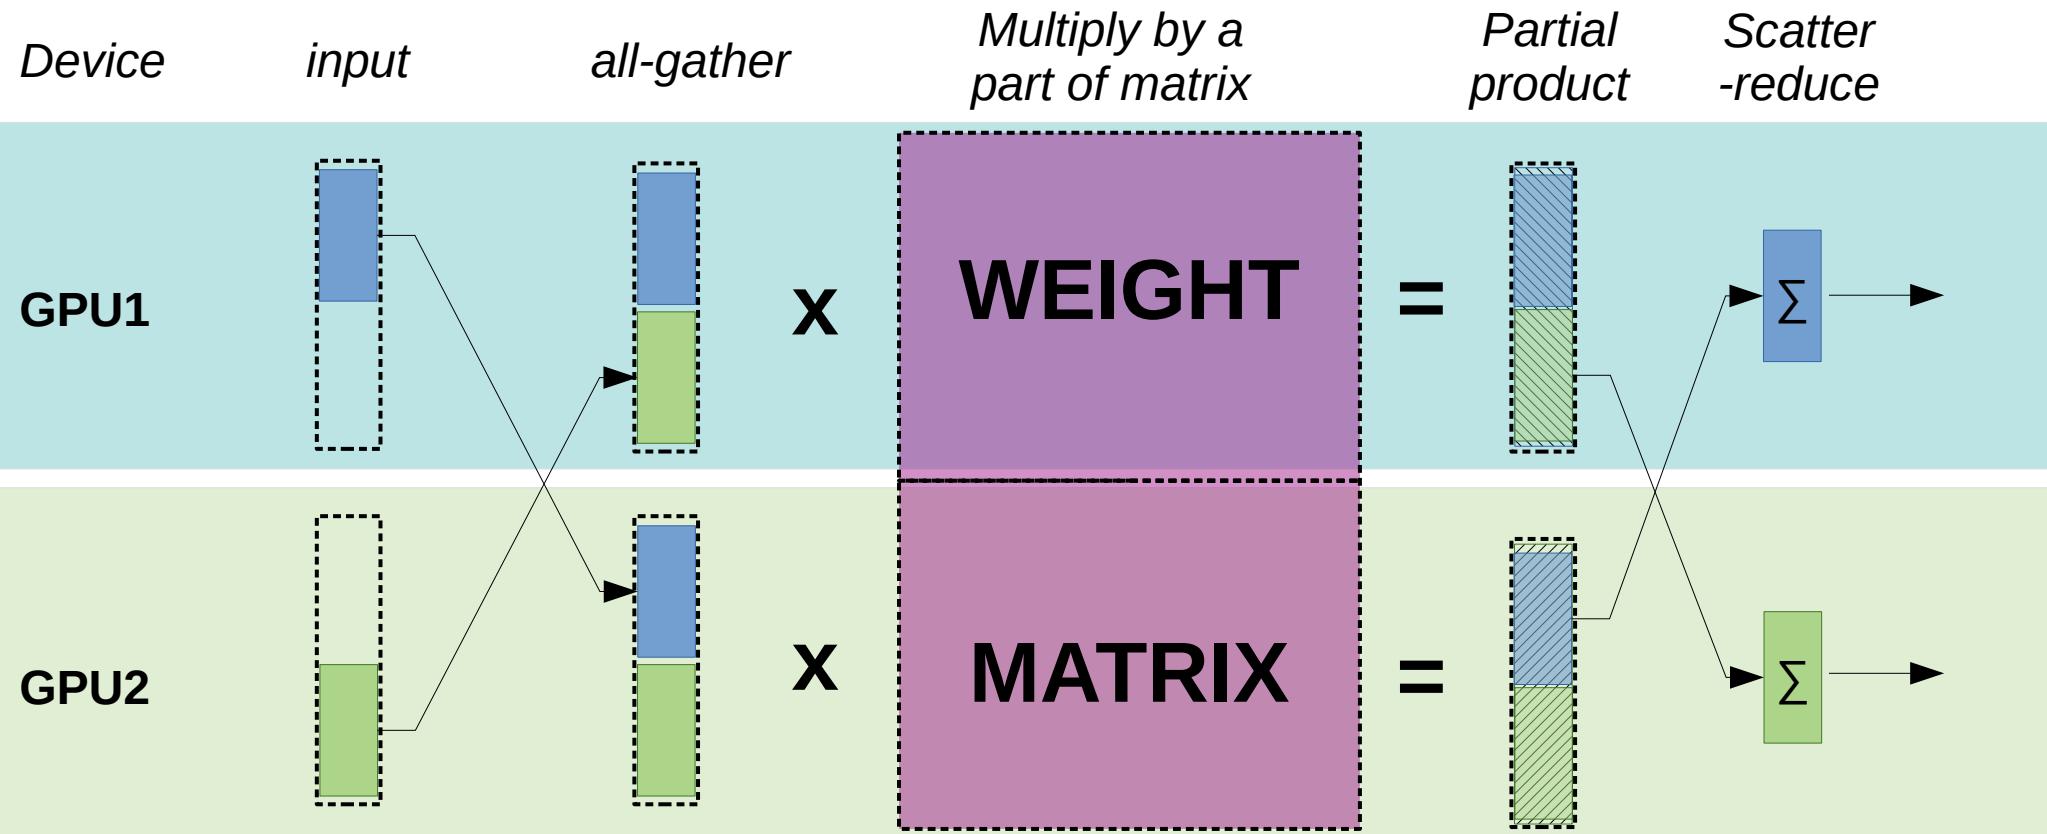
Figure 2: An illustration of the architecture of our CNN, explicitly showing the delineation of responsibilities between the two GPUs. One GPU runs the layer-parts at the top of the figure while the other runs the layer-parts at the bottom. The GPUs communicate only at certain layers. The network's input is 150,528-dimensional, and the number of neurons in the network's remaining layers is given by 253,440–186,624–64,896–64,896–43,264–4096–4096–1000.

See also: DP + TP <https://arxiv.org/abs/1404.5997>

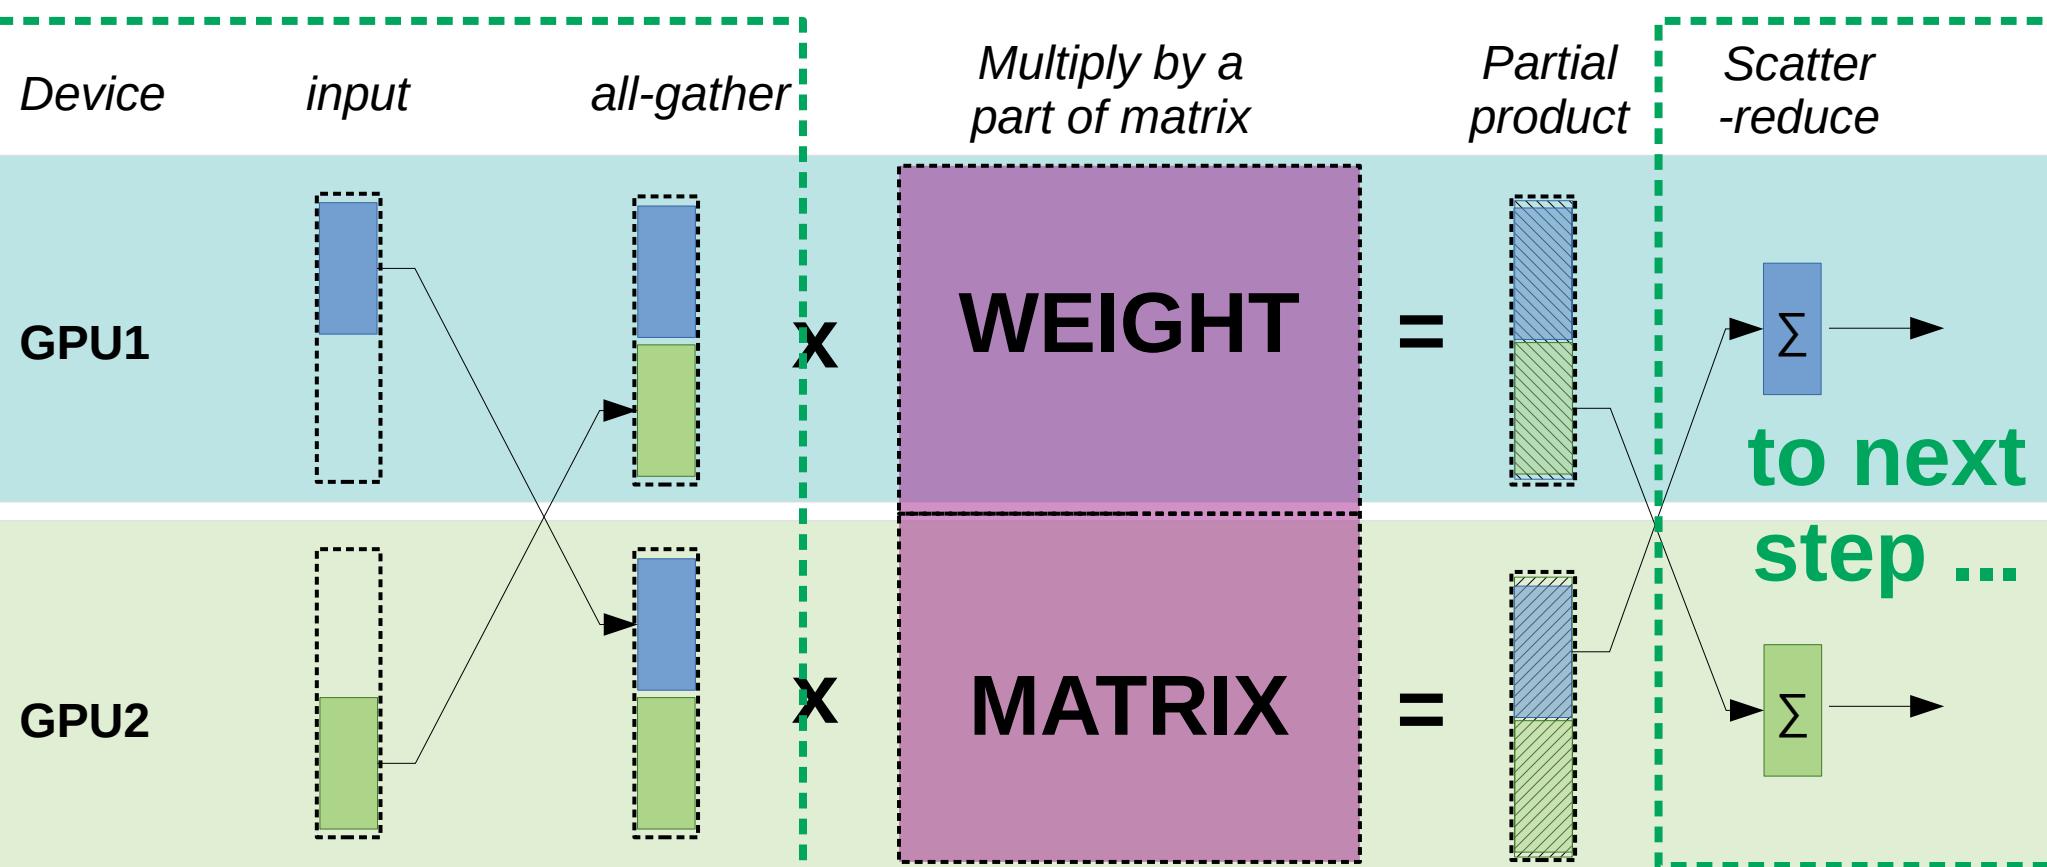
Tensor-parallel training



Q: find AllReduce op here



Q: find AllReduce op here

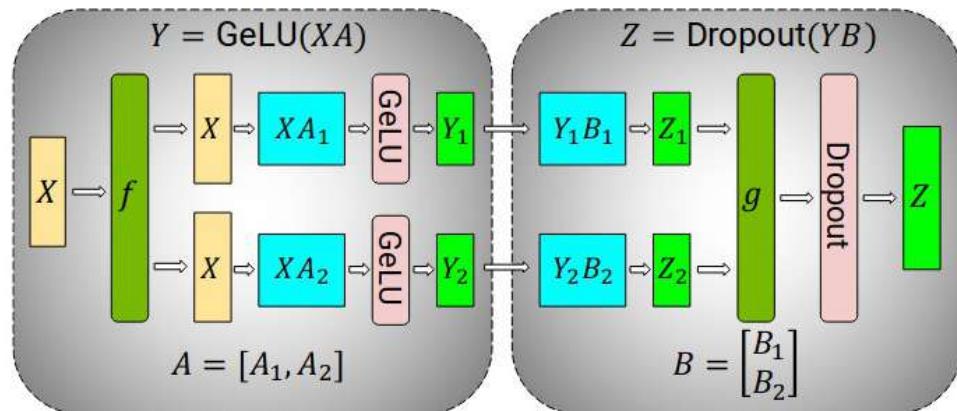


Tensor-parallel training

<https://arxiv.org/pdf/2104.04473>

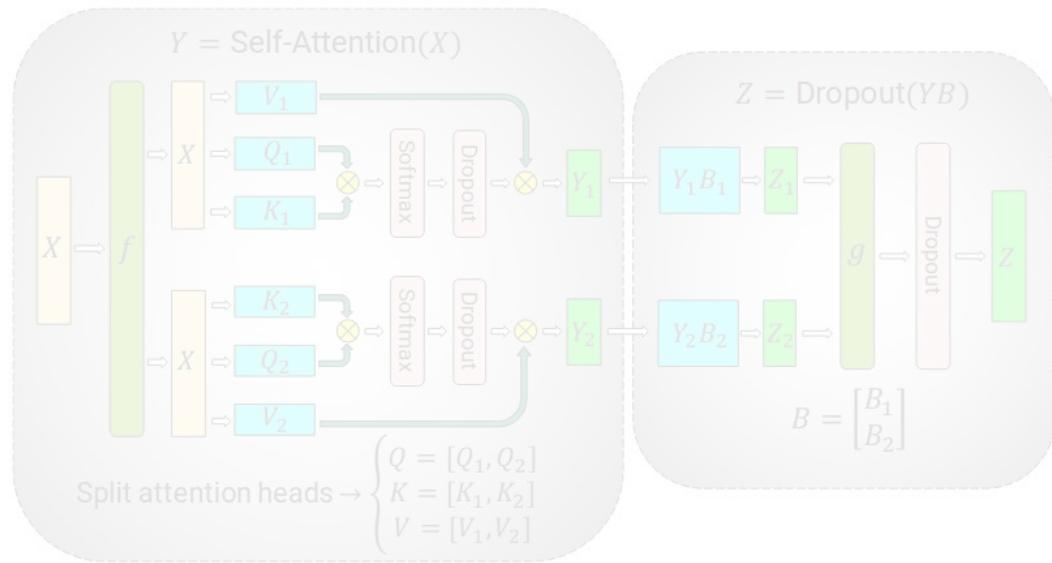
Mix and match parallelism directions to reduce synchronization

MLP: split over neurons



(a) MLP.

Attention: split over heads



(b) Self-Attention.

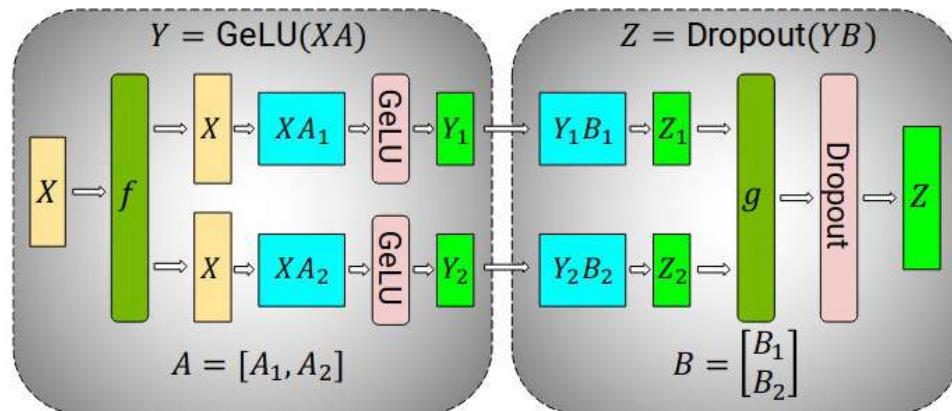
Figure 5: Blocks of transformer model partitioned with tensor model parallelism (figures borrowed from Megatron [40]). f and g are conjugate. f is the identity operator in the forward pass and all-reduce in the backward pass, while g is the reverse.

Tensor-parallel training

<https://arxiv.org/pdf/2104.04473>

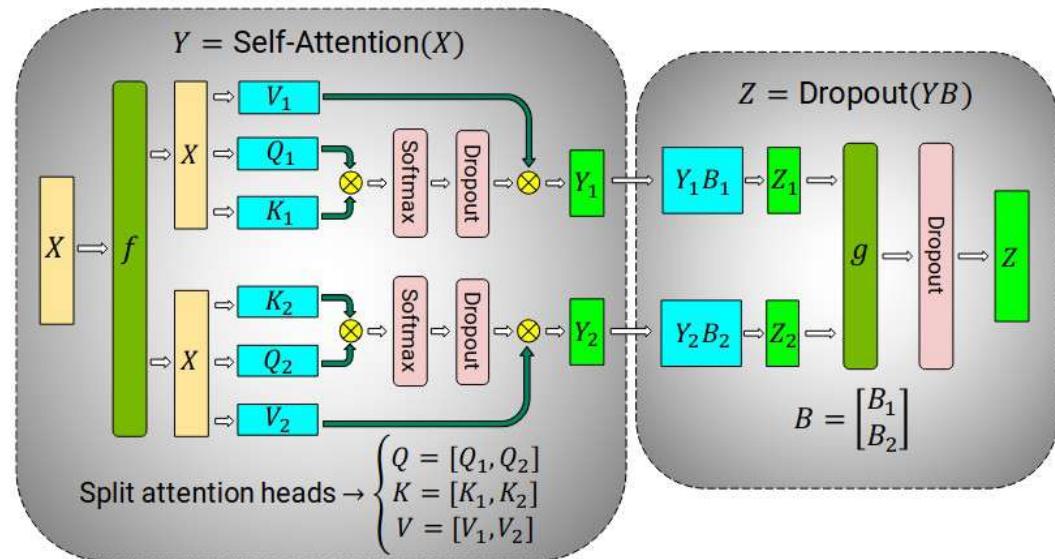
Mix and match parallelism directions to reduce synchronization

MLP: split over neurons



(a) MLP.

Attention: split over heads



(b) Self-Attention.

Figure 5: Blocks of transformer model partitioned with tensor model parallelism (figures borrowed from Megatron [40]). f and g are conjugate. f is the identity operator in the forward pass and all-reduce in the backward pass, while g is the reverse.

Sequence Parallelism

<https://arxiv.org/abs/2309.14509>

Avoid storing all activations on every device

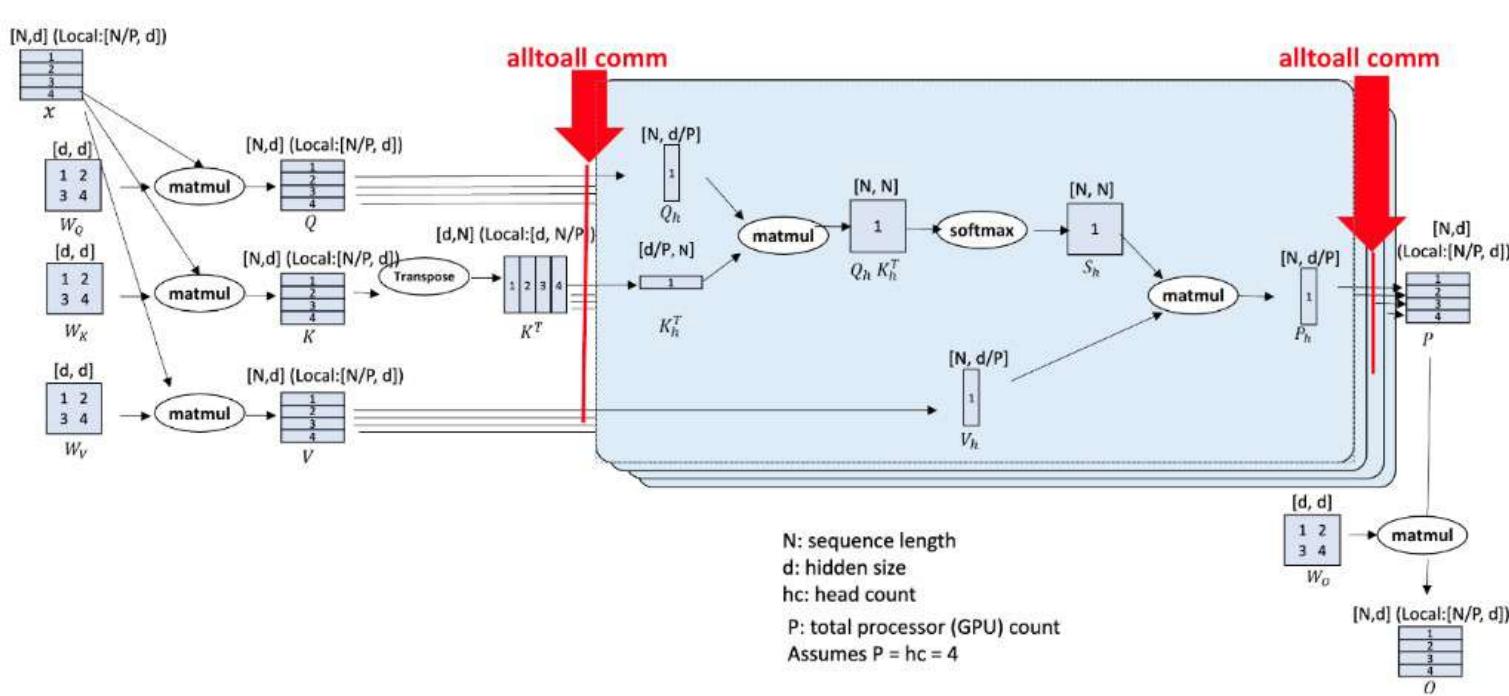


Figure 2: DeepSpeed sequence parallelism (DeepSpeed-Ulysses) design

[MOAR] Sequence Parallelism

Early mention of parallelism over sequences

<https://arxiv.org/abs/2105.05720>

DeepSpeed Ulysses – the method from previous slide

<https://arxiv.org/abs/2309.14509>

Ring Attention – compute attention dot / softmax in parallel

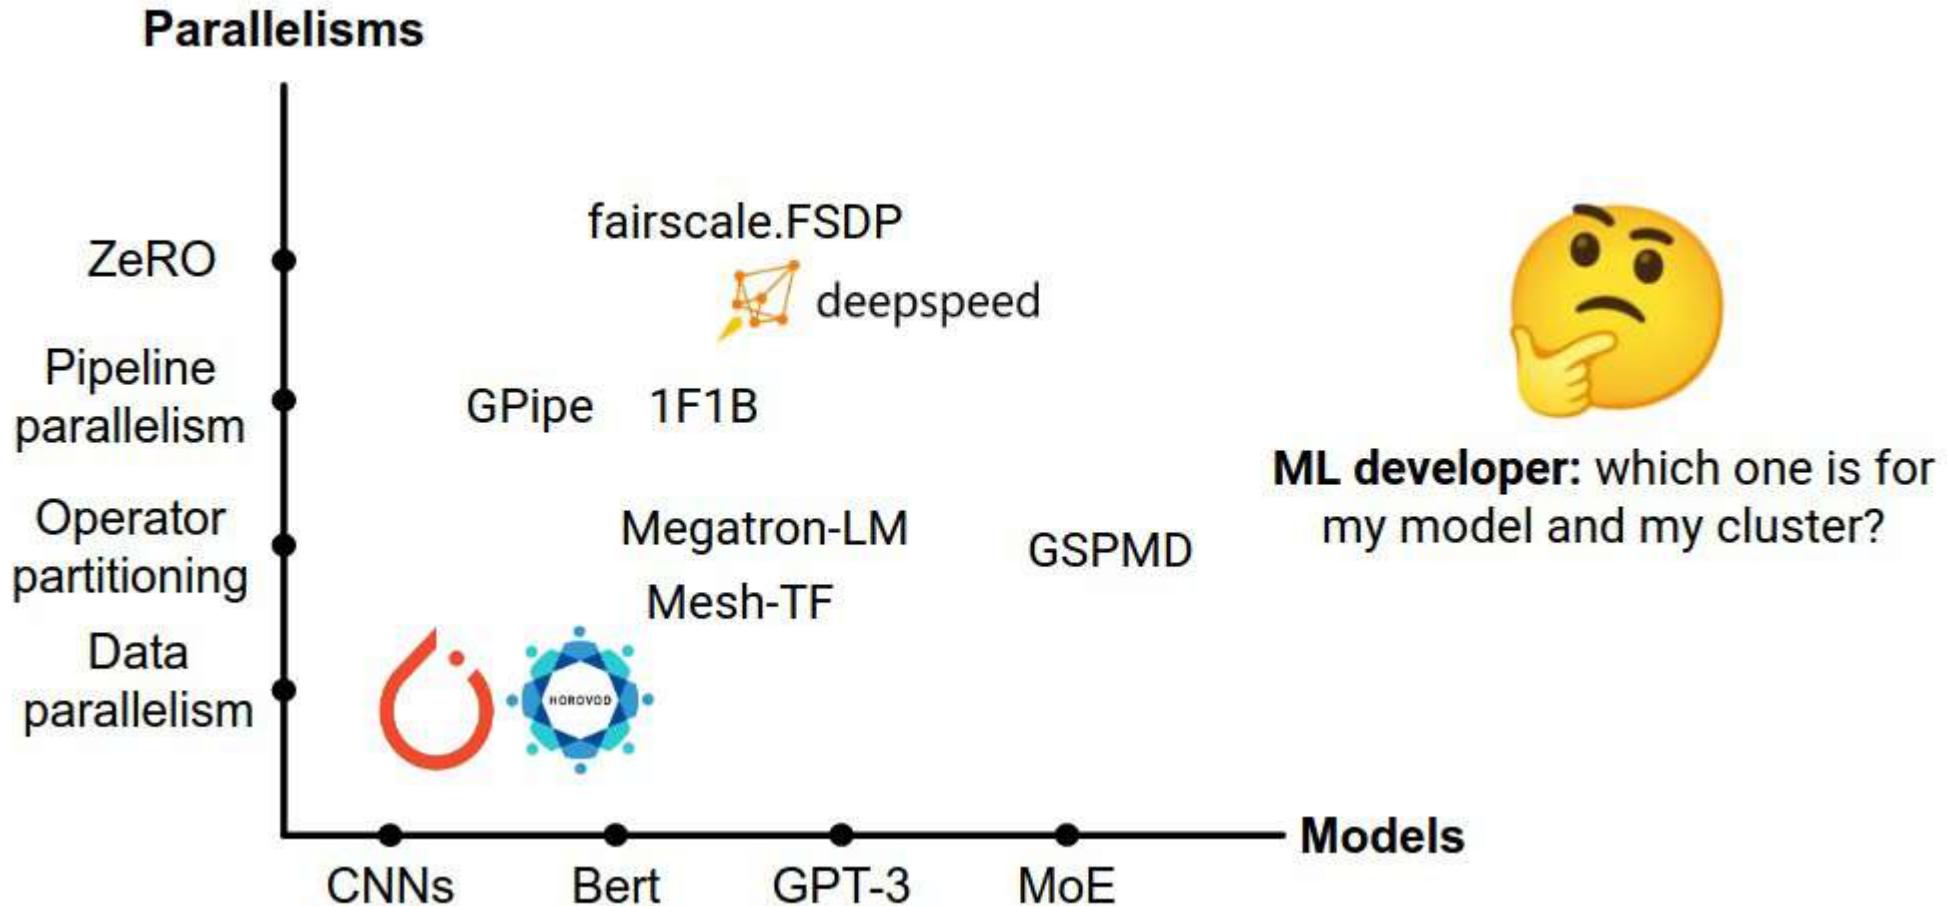
<https://arxiv.org/abs/2310.01889>

FLUX – overlap computation and communication with custom kernels

<https://arxiv.org/abs/2406.06858>

Automated parallelism

source: <https://sites.google.com/view/icml-2022-big-model>



Automated parallelism

source: <https://sites.google.com/view/icml-2022-big-model>

Classic view

Data parallelism

Model parallelism

New view (this tutorial)

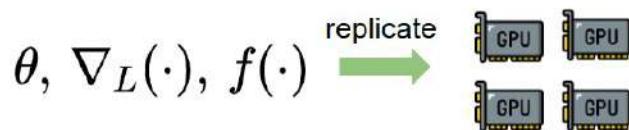
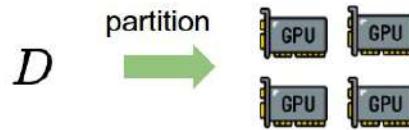
Inter-op parallelism

Intra-op parallelism

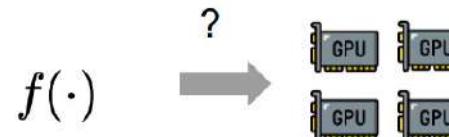
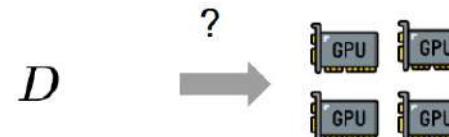
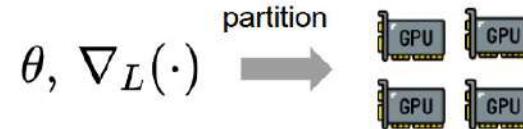
Automated parallelism

source: <https://sites.google.com/view/icml-2022-big-models>

Data parallelism



Model parallelism



The diagram shows a sequence of operations: **data** → **model** (e.g., CNN, GPT, etc.) → **output** → **weight update** (e.g., SGD, adam, etc.) → **parameter** → **theta^(t+1)**. The **theta^(t+1)** is shown as a bracketed expression: $\theta^{(t+1)} = f(\theta^{(t)}, \nabla_L(\theta^{(t)}, D^{(t)}))$.

Automated parallelism

source: <https://sites.google.com/view/icml-2022-big-model>

Data and model parallelism

- Two pillars: **data** and **model**.
- ✓ “Data parallelism” is general and precise.
- ? “Model parallelism” is vague.
- ? The view creates ambiguity for methods that neither partitions data nor the model computation.

New: Inter-op and Intra-op parallelism.

- Two pillars: **computational graph** and **device cluster**
- ✓ This view is based on their computing characteristics.
- ✓ This view facilitates the development of new parallelism methods.

Automated parallelism

source: <https://sites.google.com/view/icml-2022-big-model>

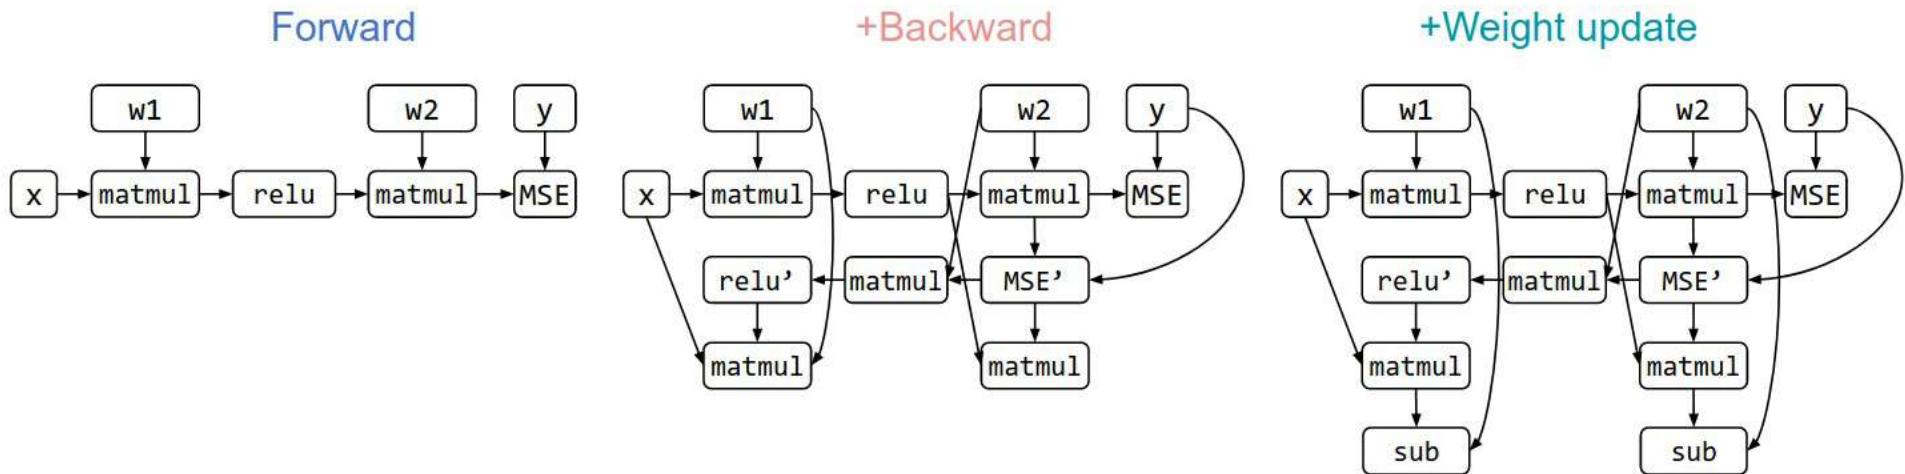
$$\theta^{(t+1)} = f(\theta^{(t)}, \nabla_L(\theta^{(t)}, D^{(t)}))$$

$$L = \text{MSE}(w_2 \cdot \text{ReLU}(w_1 x), y) \quad \theta = \{w_1, w_2\}, D = \{(x, y)\}$$

$$f(\theta, \nabla_L) = \theta - \nabla_L$$

□ Operator / its output tensor

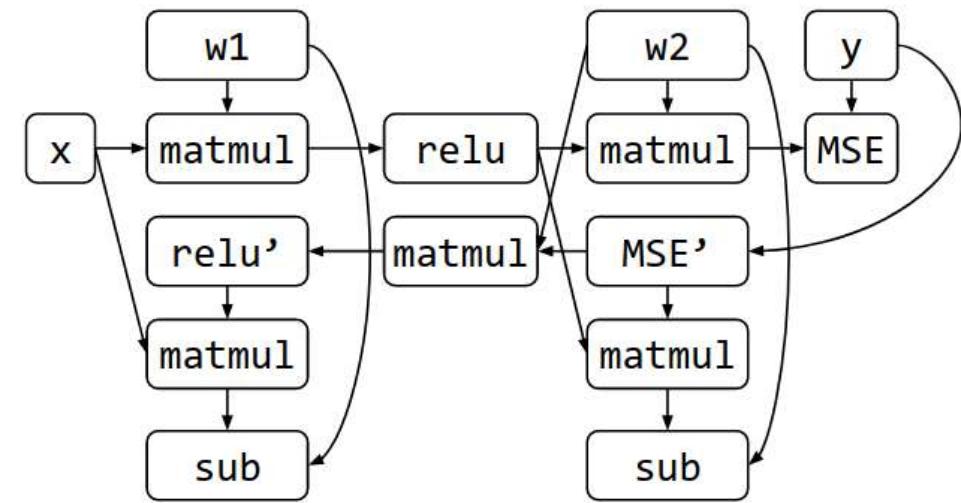
→ Data flowing direction



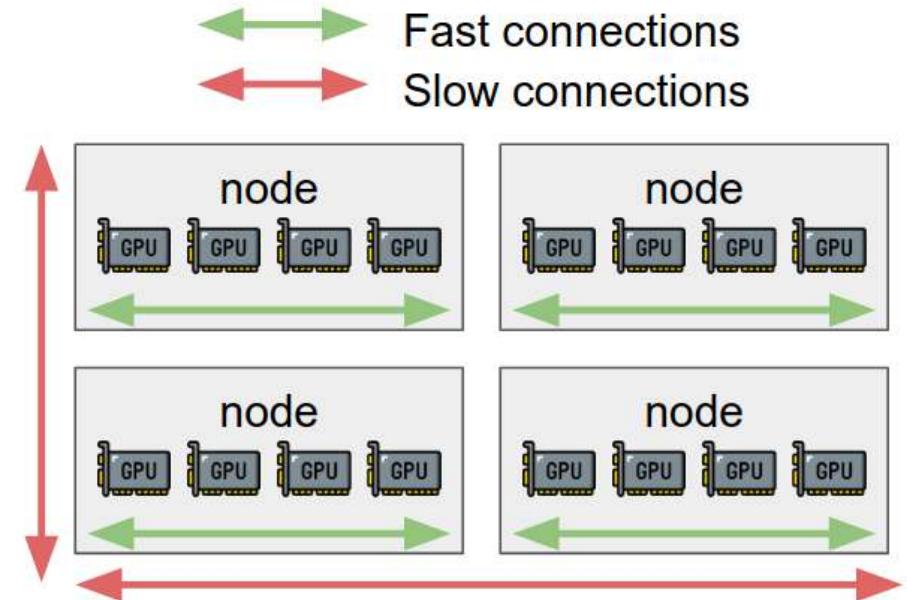
Automated parallelism

source: <https://sites.google.com/view/icml-2022-big-model>

Compute graph



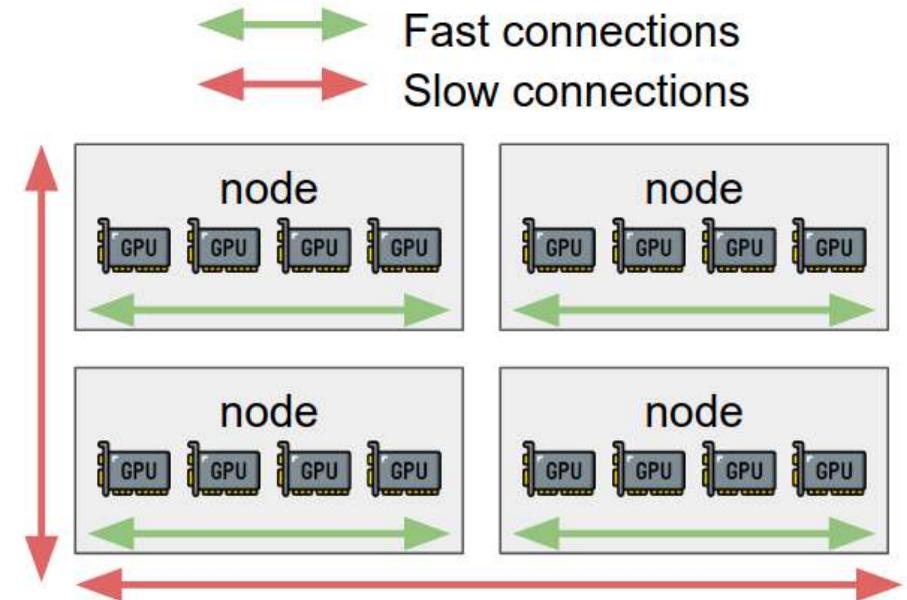
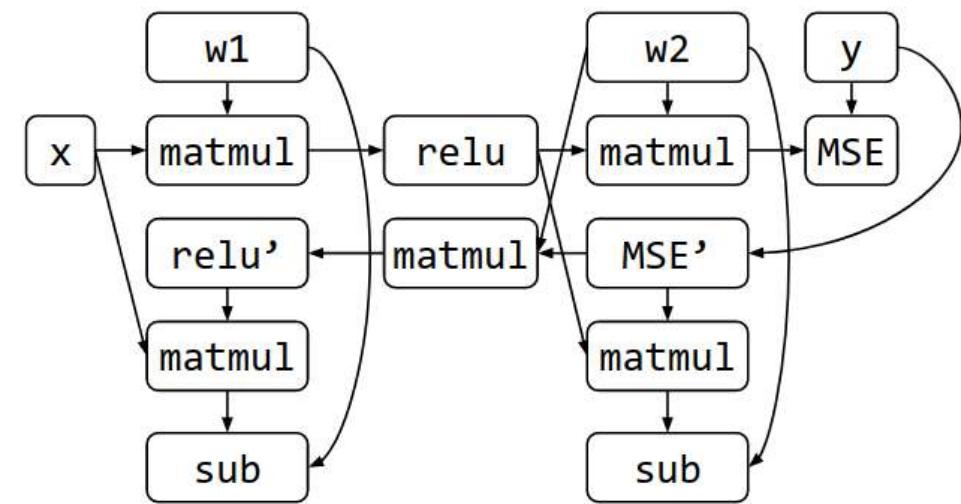
Device cluster



Automated parallelism

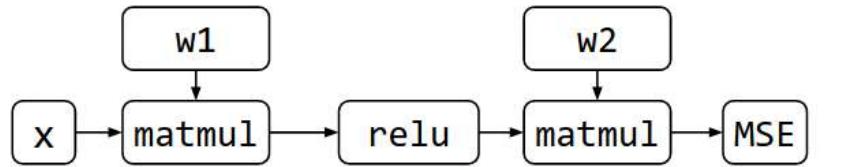
source: <https://sites.google.com/view/icml-2022-big-model>

Q: How to partition the graph on the device cluster?



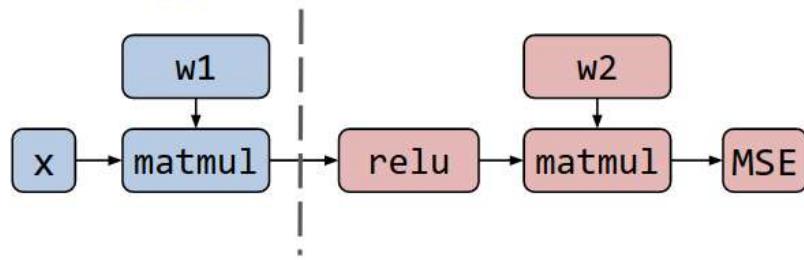
Automated parallelism

source: <https://sites.google.com/view/icml-2022-big-model>

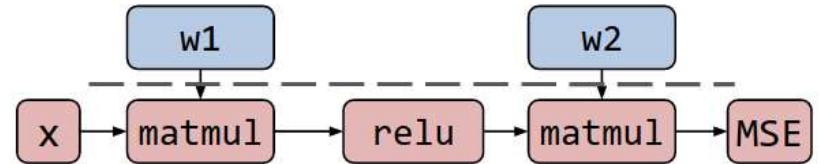


Device 1
Device 2

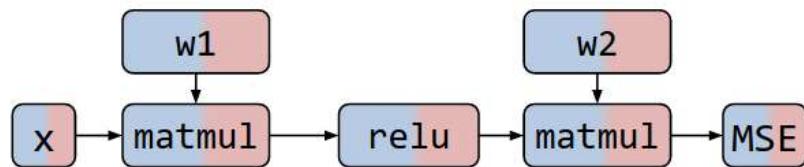
Strategy 1



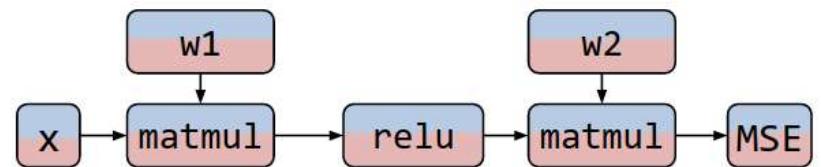
Strategy 2



Strategy 3

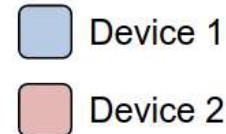
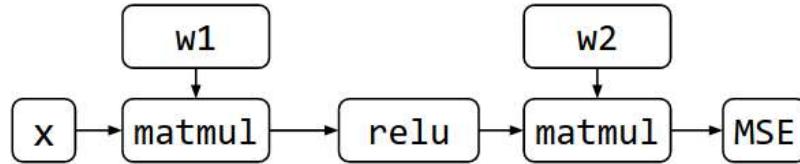


Strategy 4



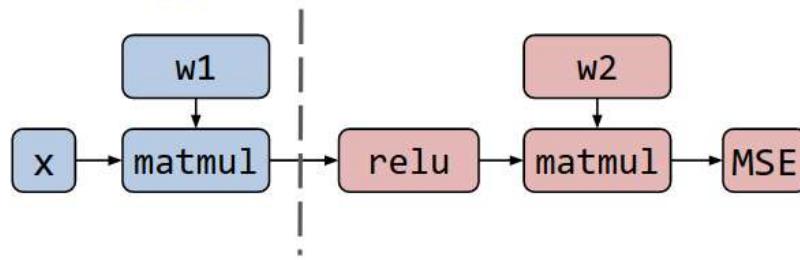
Automated parallelism

source: <https://sites.google.com/view/icml-2022-big-model>

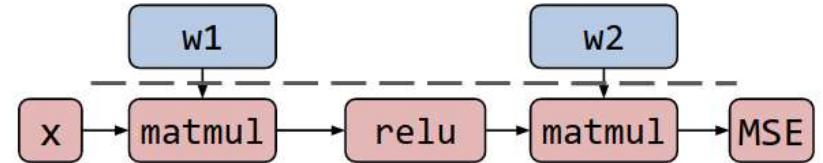


Q: have you seen S1/2/3/4 before?

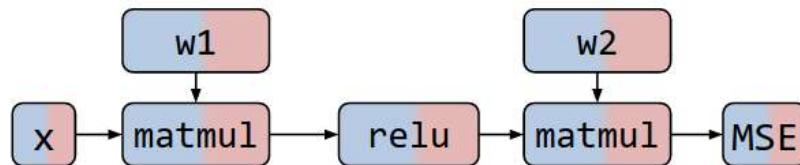
Strategy 1



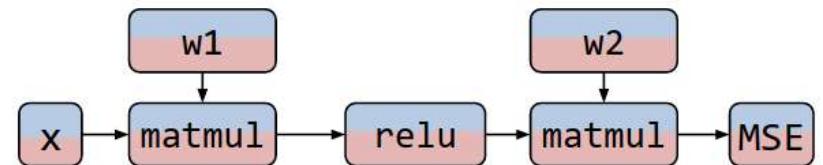
Strategy 2



Strategy 3

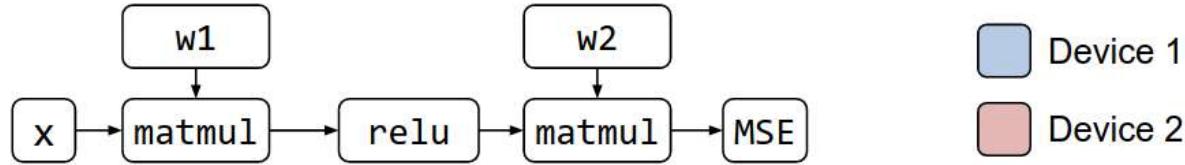


Strategy 4

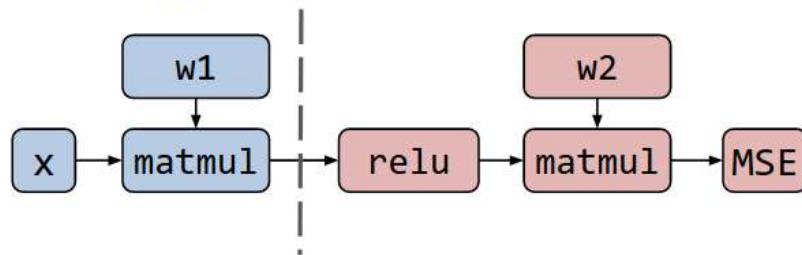


Automated parallelism

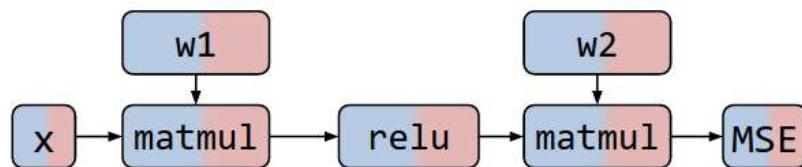
source: <https://sites.google.com/view/icml-2022-big-model>



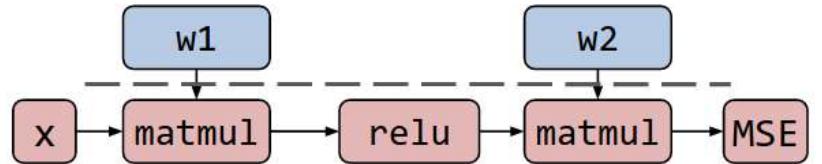
Pipeline MP



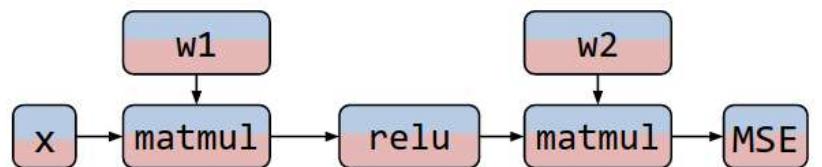
Tensor-parallel v1



DP with offloading or PS

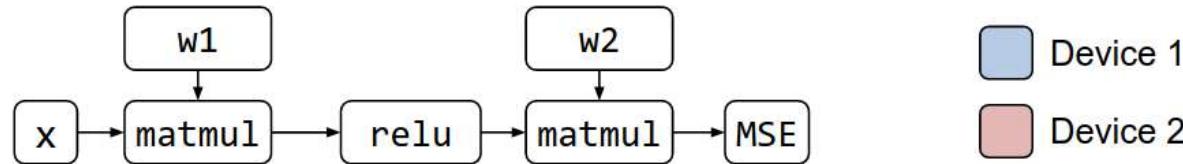


Tensor-parallel v2

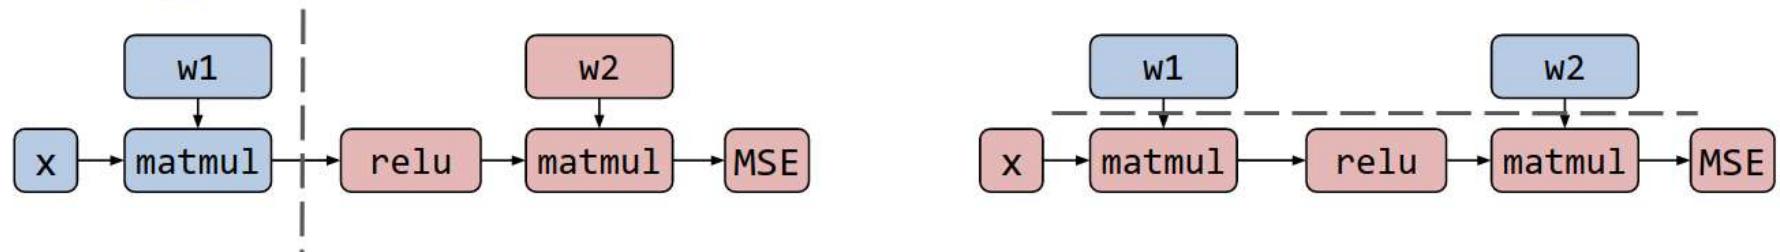


Automated parallelism

source: <https://sites.google.com/view/icml-2022-big-model>



~~Pipeline MP~~ **Inter-op parallelism** ~~DP with offloading or PS~~

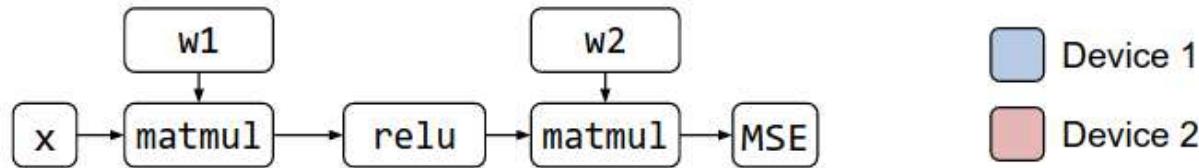


~~Tensor-parallel v1~~ **Intra-op parallelism** ~~Tensor-parallel v2~~



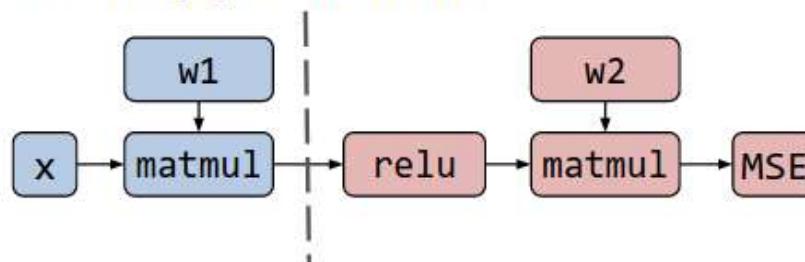
Automated parallelism

source: <https://sites.google.com/view/icml-2022-big-model>



Device 1
Device 2

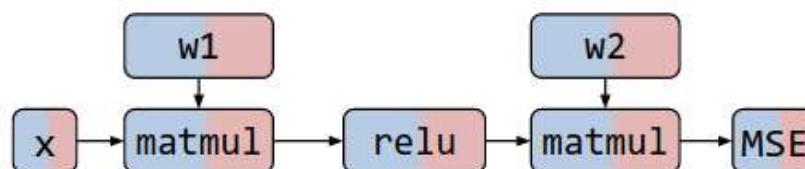
Inter-op parallelism



Trade-off

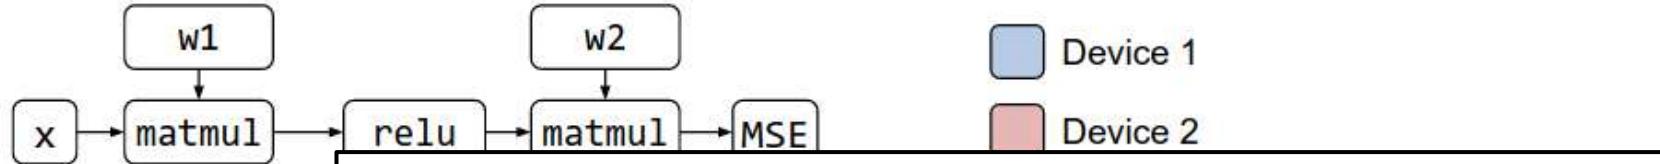
Inter-operator Parallelism	Intra-operator Parallelism
Communication	Less
Device Idle Time	More

Intra-op parallelism

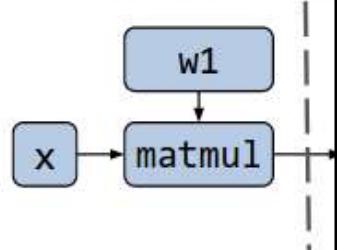


Automated parallelism

source: <https://sites.google.com/view/icml-2022-big-model>

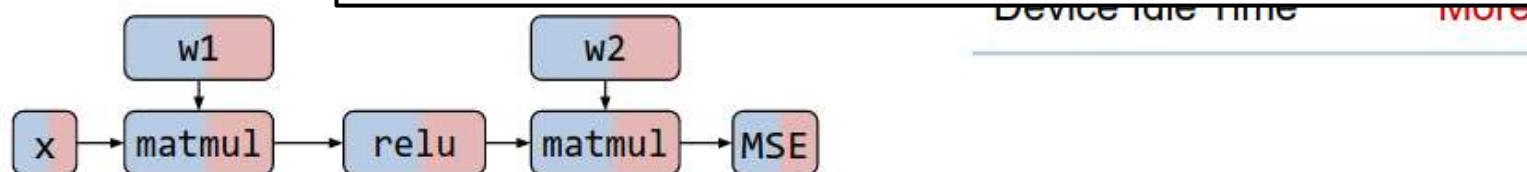


Inter-op parallel



Q: how do we find the best strategy
for partitioning the graph?

Intra-op parallel



Intra-operator
Parallelism

More

Less

RL-based partitioning

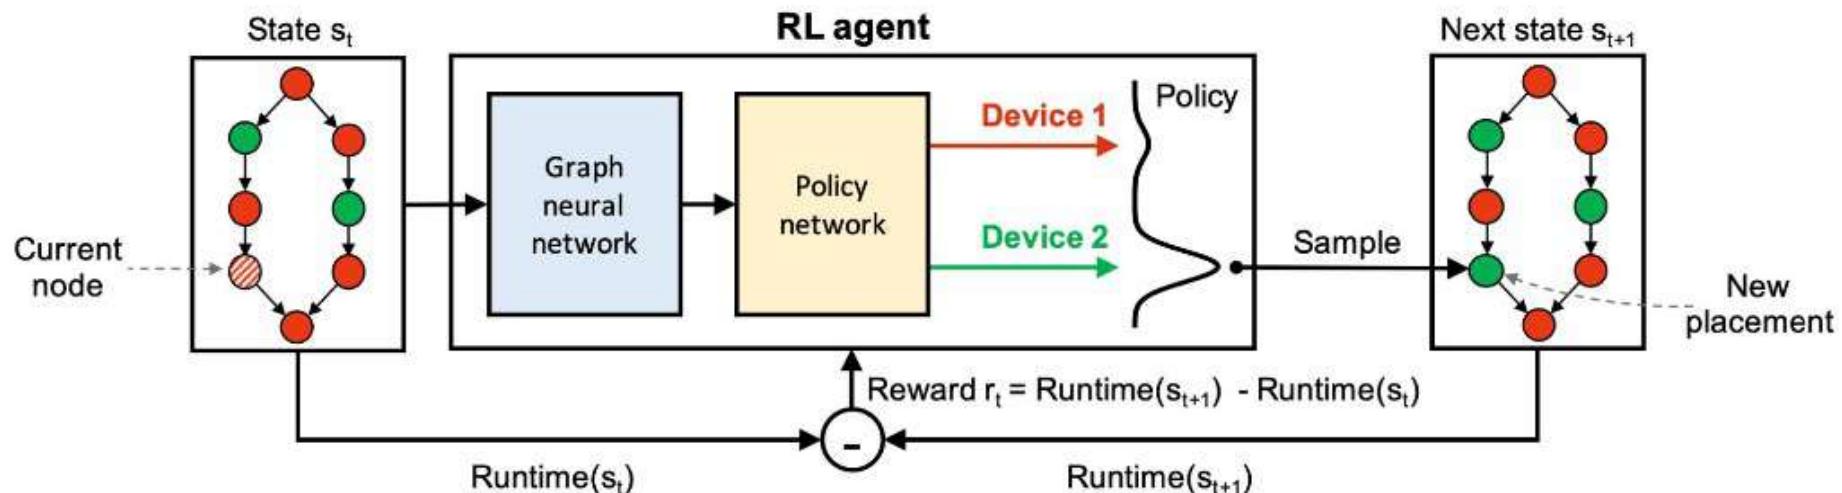
<https://people.csail.mit.edu/hongzi/content/publications/placeto-neurips19.pdf>

State: Device assignment plan for a computational graph.

Action: Modify the device assignment of a node.

Reward: Latency difference between the new and old placements.

Trained with **policy gradient** algorithm.



Optimization-based partitioning

<https://arxiv.org/abs/2006.16423>

Integer Linear Programming:

Variable: Decision variable vector for each operator, representing device assignment.

Minimize: Maximum finishing time of all operators.

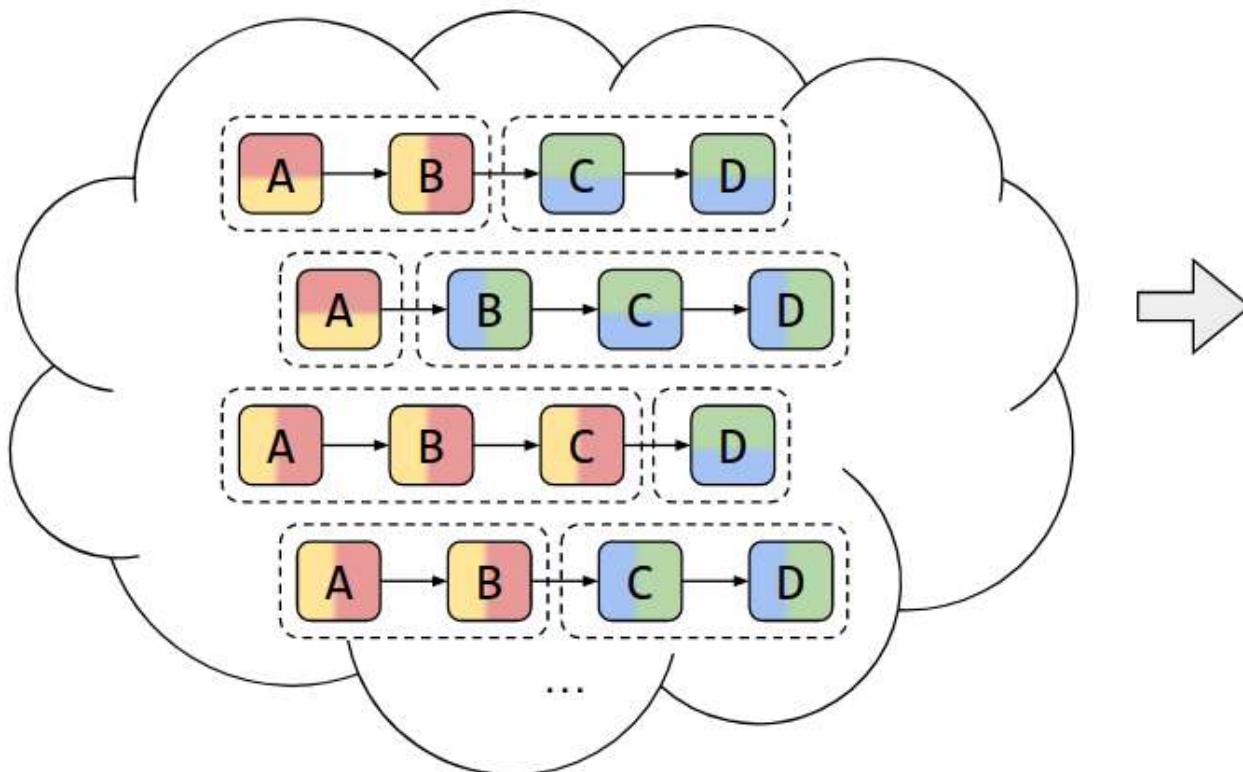
Constraint: Execution dependency & memory capacity of each device.

$$\begin{aligned} \min \quad & \text{TotalLatency} \\ \text{s.t.} \quad & \sum_{i=0}^k x_{vi} = 1 \\ & \text{subgraph } \{v \in V : x_{vi} = 1\} \text{ is contiguous} \\ & M \geq \sum_v m_v \cdot x_{vi} \\ & \text{CommIn}_{ui} \geq x_{vi} - x_{ui} \\ & \text{CommOut}_{ui} \geq x_{ui} - x_{vi} \\ & \text{TotalLatency} \geq \text{Latency}_v \\ & \text{SubgraphStart}_i \geq \text{Latency}_v \cdot \text{CommIn}_{vi} \\ & \text{SubgraphFinish}_i = \text{SubgraphStart}_i + \sum_v \text{CommIn}_{vi} \cdot c_v \\ & \quad + \sum_v x_{vi} \cdot p_v^{\text{acc}} + \sum_v \text{CommOut}_{vi} \cdot c_v \\ & \text{Latency}_v \geq x_{v0} \cdot p_v^{\text{cpu}} \\ & \text{Latency}_v \geq x_{v0} \cdot p_v^{\text{cpu}} + \text{Latency}_u \\ & \text{Latency}_v \geq x_{vi} \cdot \text{SubgraphFinish}_i \\ & x_{vi} \in \{0, 1\} \end{aligned}$$

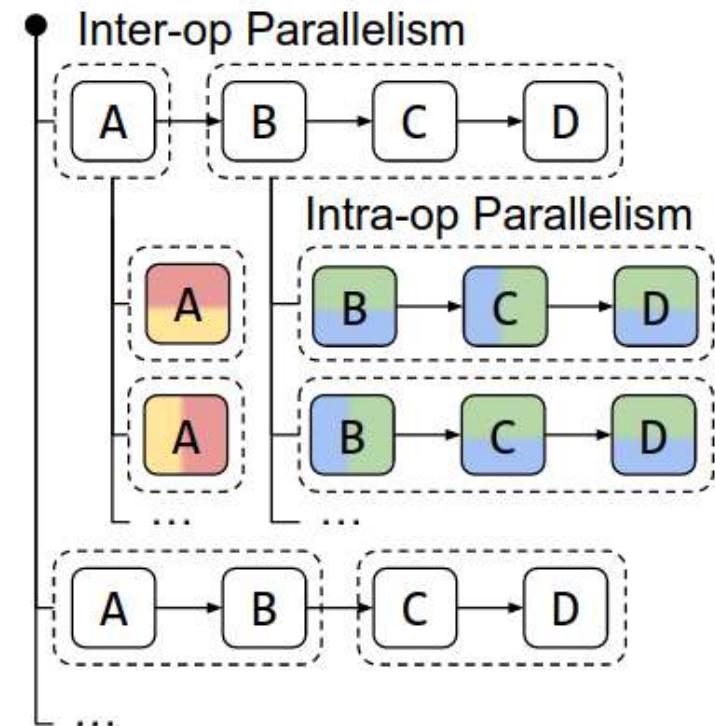
Alpa: optimization-based + reduced search space

<https://arxiv.org/abs/2201.12023>

Whole Search Space

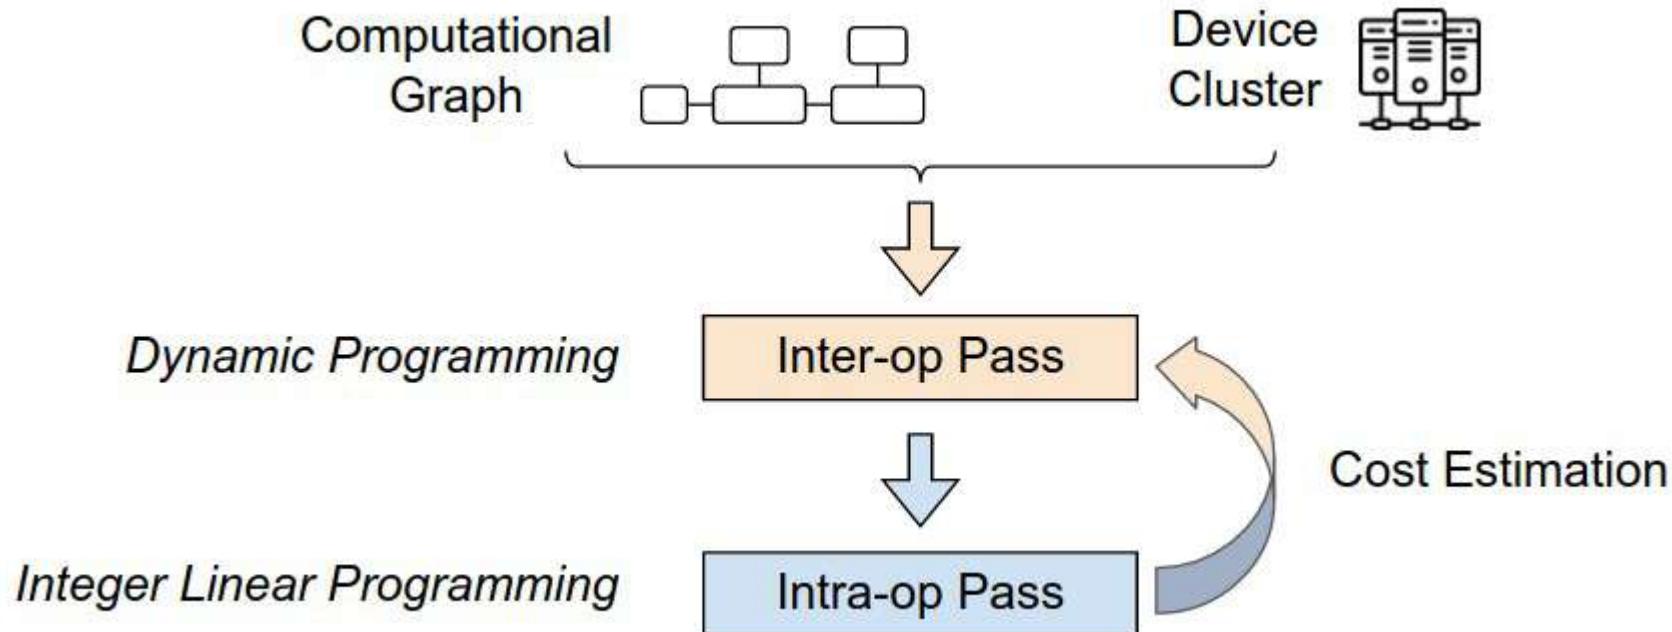


Alpa Hierarchical Space



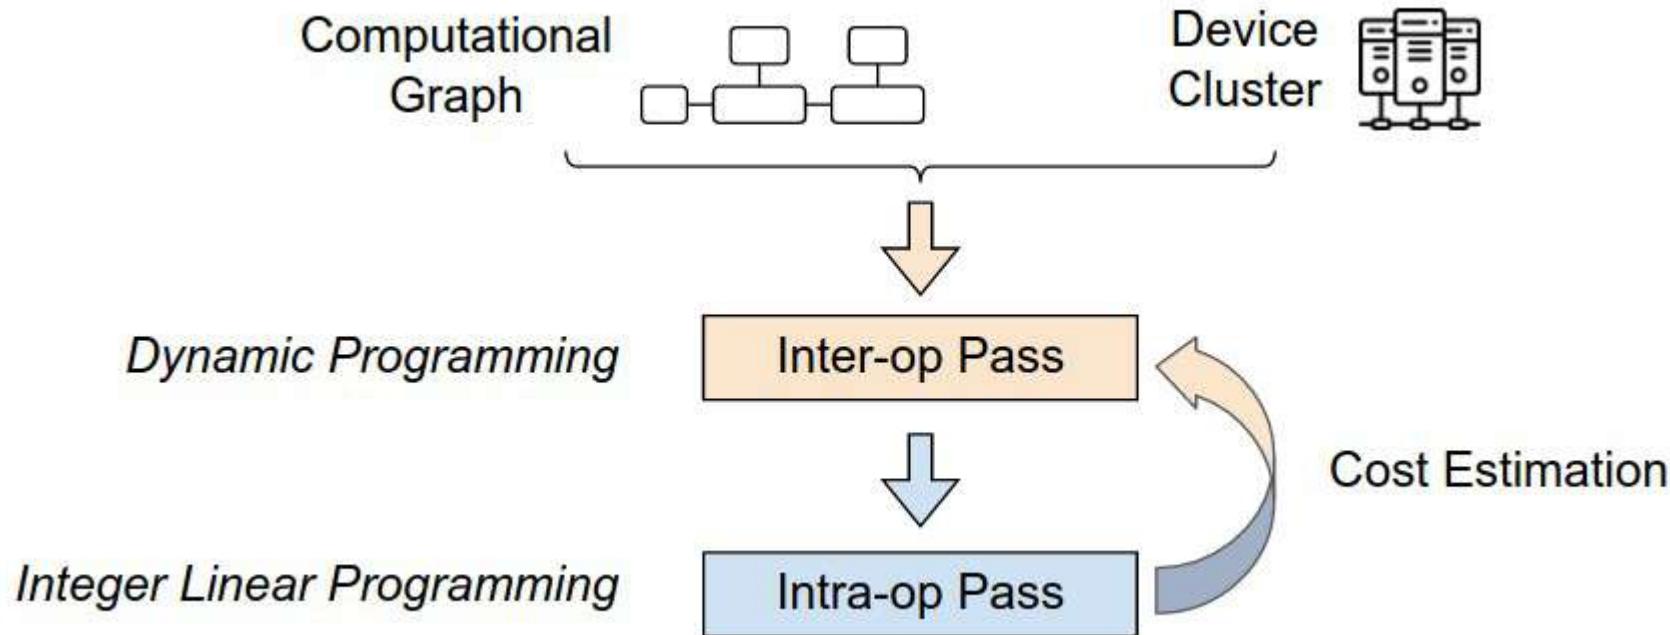
Alpa: optimization-based + reduced search space

<https://arxiv.org/abs/2201.12023>



Alpa: optimization-based + reduced search space

<https://arxiv.org/abs/2201.12023>



More details of each pass:

<https://sites.google.com/view/icml-2022-big-model>

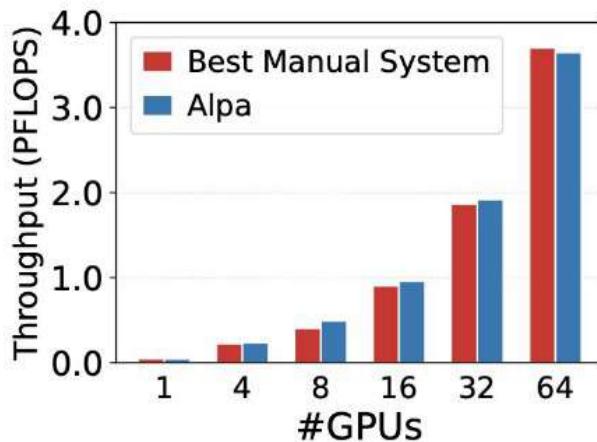
Alpa: optimization-based + reduced search space

<https://arxiv.org/abs/2201.12023>

Not the first algorithm for auto-parallelism...
but the first one that is usable* (* - most of the time)

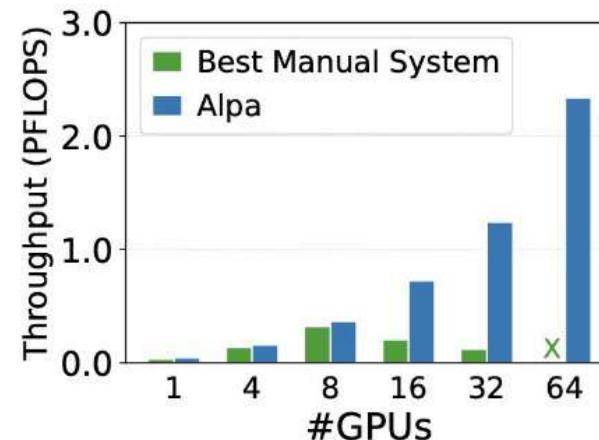
(benchmarks on AWS V100)

GPT (up to 39B)



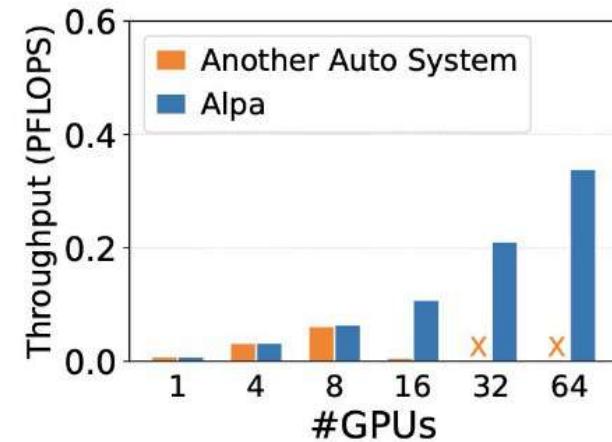
Match specialized
manual systems.

GShard MoE (up to 70B)



Outperform the manual
baseline by up to 8x.

Wide-ResNet (up to 13B)



Generalize to models
without manual plans.

Alpa: optimization-based + reduced search space

<https://arxiv.org/abs/2201.12023>

Not the first algorithm for auto-parallelism...
but the first one that is usable* (* - most of the time)

```
# Define the training step. The body of this function is the same as the
# ``train_step`` above. The only difference is to decorate it with
# ``alpa.parallelize``.

@alpa.parallelize auto best strategy
def alpa_train_step(state, batch):
    def loss_func(params):
        out = state.apply_fn(params, batch["x"])
        loss = jnp.mean((out - batch["y"])**2)
        return loss works in jax

    grads = jax.grad(loss_func)(state.params)
    new_state = state.apply_gradients(grads=grads)
    return new_state

# Test correctness
actual_state = alpa_train_step(state, batch)
assert_allclose(expected_state.params, actual_state.params, atol=5e-3)
```

Alpa: optimization-based + reduced search space

<https://arxiv.org/abs/2201.12023>

Not the first algorithm for auto-parallelism...

but the first one that is usable* (* - most of the time)

```
# Define training graph
# ``train`` function
# ``alpa_train_step`` function
@alpa.pipeline
def alpa_train_step(state, batch):
    def forward_fn(state, batch):
        # Compute loss
        loss = state.loss_fn(batch)
        return loss
```

Alpa was deprecated in 2024, but successors exist

Jax: use `pjit/xmap` with improved XLA optimizations

**PyTorch: <https://docs.pytorch.org/xla/master/spmd.html>
... but for standard models, DeepSpeed is often enough.**

```
graph
new_state = state
ret = loss
```

```
# Test correctness
actual_state = alpa_train_step(state, batch)
assert_allclose(expected_state.params, actual_state.params, atol=5e-3)
```

</part 2>

- + model larger than GPU
- + faster for small
 - * typical size: 2-8 gpus
- model partitioning is tricky
 - tensor parallelism is easier, but requires ultra low latency
- latency is critical, go buy nvlink
 - except for *PipeDream*
- *often combined with gradient checkpointing*

Tutorials:

- Simple pipelining in PyTorch – tinyurl.com/pytorch-pipelining
- Distributed model-parallel with torch RPC - <https://tinyurl.com/torch-rpc>
- Minimalistic tensor parallelism [pip install tensor_parallel](#)

</part 2>

- + model larger than GPU
- + faster for small
 - * typical size: 2-8 gpus
- model partitioning is tricky
 - tensor parallelism is easier, but requires ultra low latency
- latency is critical, go buy nvlink
 - except for *PipeDream*
- *often combined with gradient checkpointing*

Tutorials:

- Simple pipelining in PyTorch – tinyurl.com/pytorch-pipelining
- Distributed model-parallel with torch RPC - <https://tinyurl.com/torch-rpc>
- Automatic tensor parallelism `pip install tensor_parallel`

Q: what if you have 1024 GPUs, but the model fits on 8?

</part 2>

- + model larger than GPU
- + faster for small
 - * typical size: 2-8 gpus
- model partitioning is tricky
 - tensor parallelism is easier, but requires ultra low latency
- latency is critical, go buy nvlink
 - except for *PipeDream*
- often combined with gradient checkpointing

Tutorials:

- Simple pipelining in PyTorch – tinyurl.com/pytorch-pipelining
- Distributed model-parallel with torch RPC - <https://tinyurl.com/torch-rpc>
- Automatic tensor parallelism `pip install tensor_parallel`

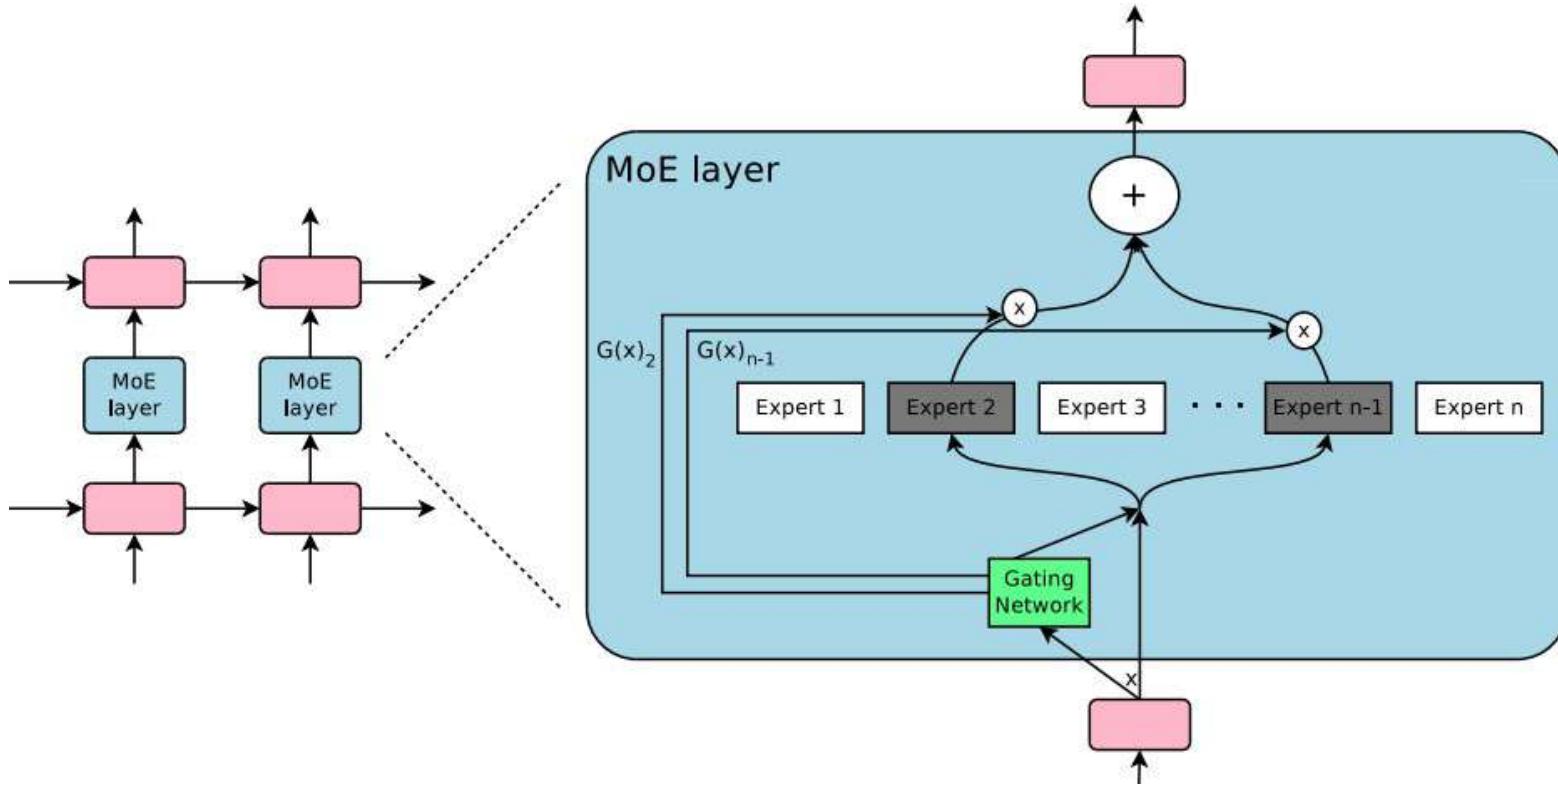
Large-scale training: combine model- and data-parallel

So far we've been trying to partition for existing models...

Perhaps there are models that are easier to partition?

Expert Parallelism

Sparingly gated MoE: <https://arxiv.org/pdf/1701.06538.pdf>

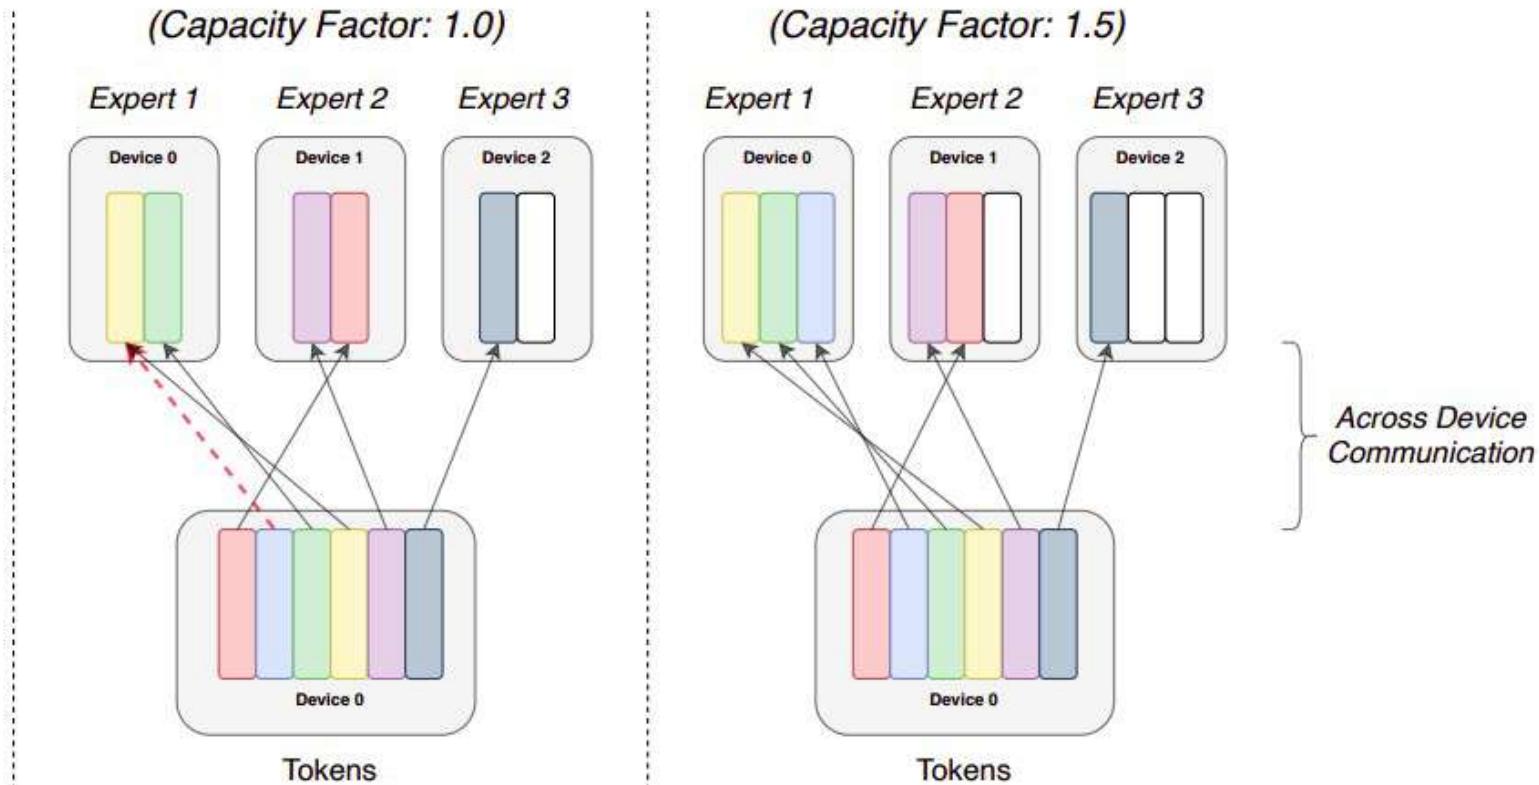


MoE Variant: Switch Transformer

Switch: <https://arxiv.org/pdf/2101.03961.pdf>

Terminology

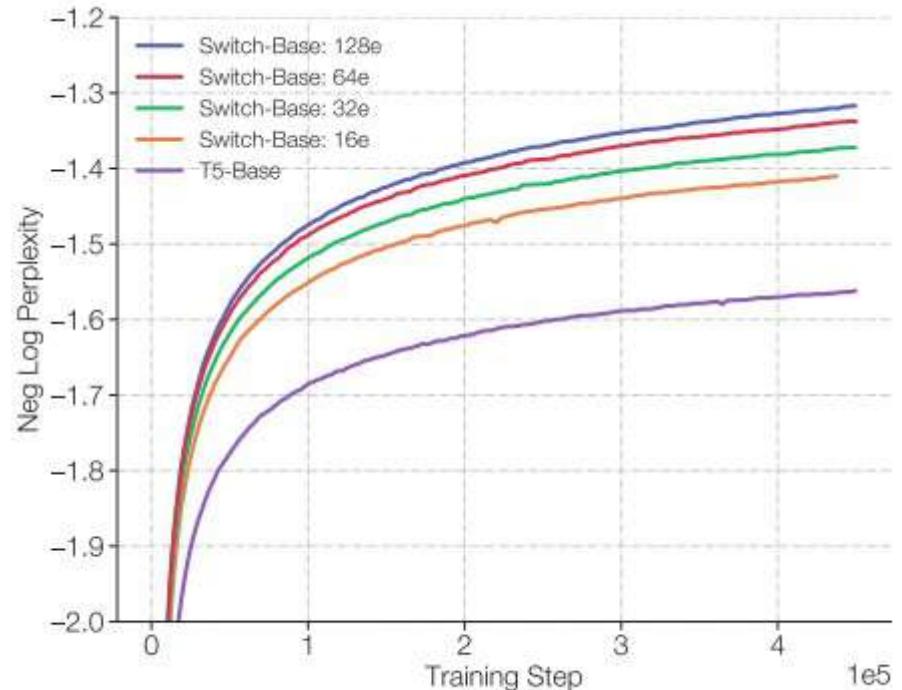
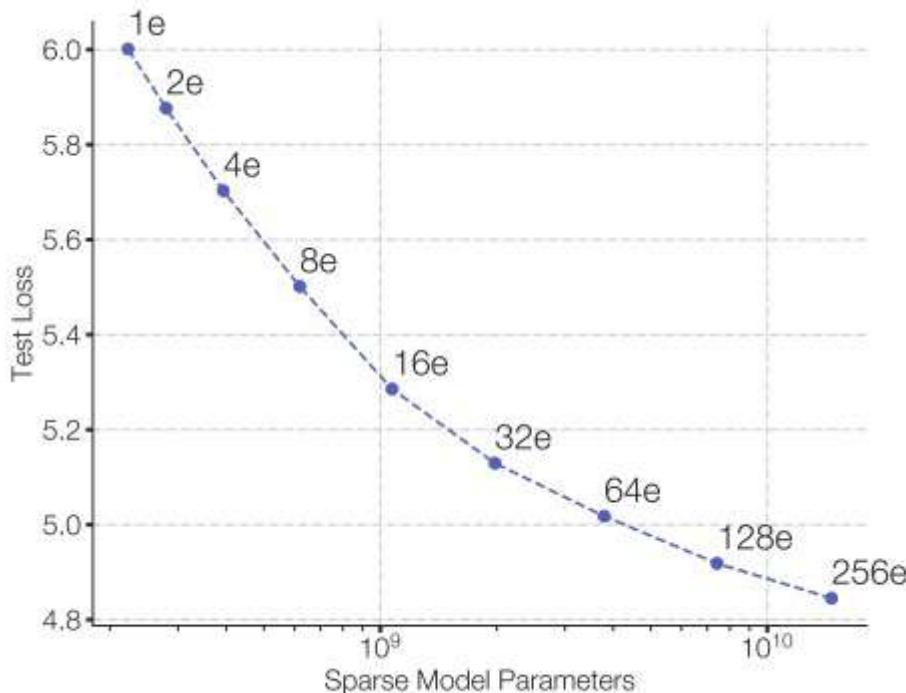
- **Experts:** Split across devices, each having their own unique parameters. Perform standard feed-forward computation.
- **Expert Capacity:** Batch size of each expert. Calculated as
$$(\text{tokens_per_batch} / \text{num_experts}) * \text{capacity_factor}$$
- **Capacity Factor:** Used when calculating expert capacity. Expert capacity allows more buffer to help mitigate token overflow during routing.



MoE Variant: Switch Transformer

Switch: <https://arxiv.org/pdf/2101.03961.pdf>

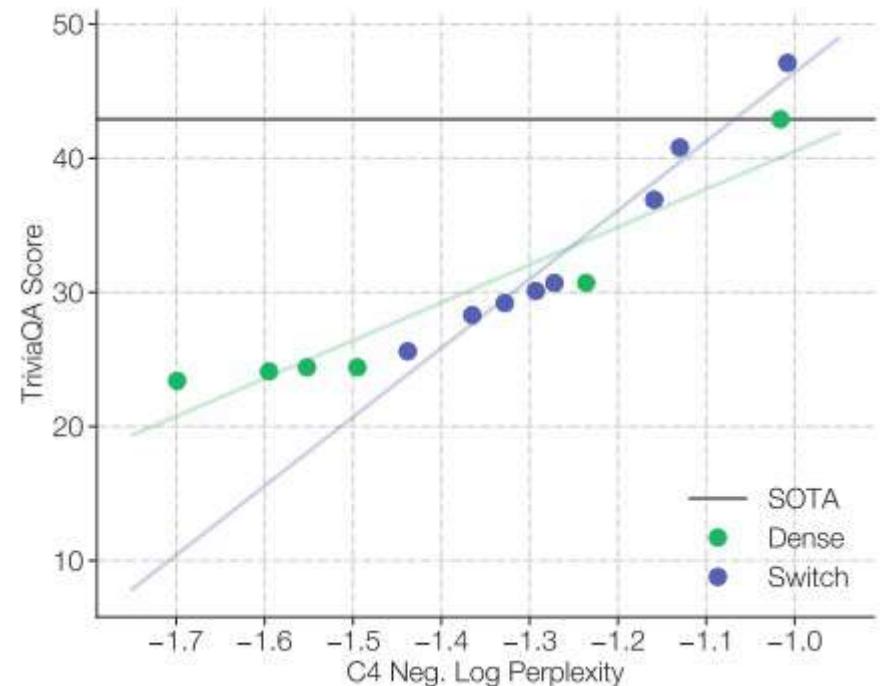
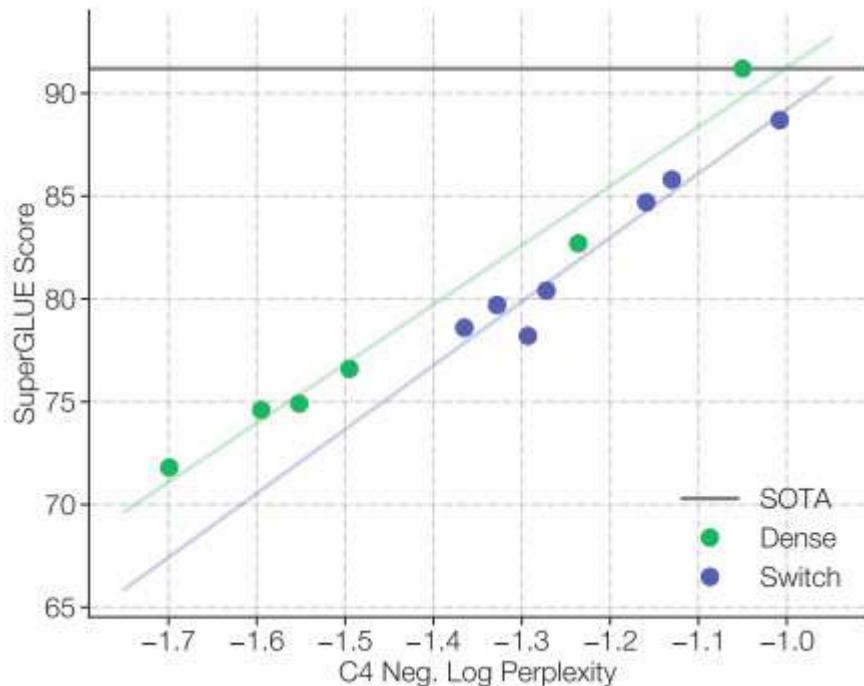
MLM pre-training objective [BERT-like]



MoE Variant: Switch Transformer

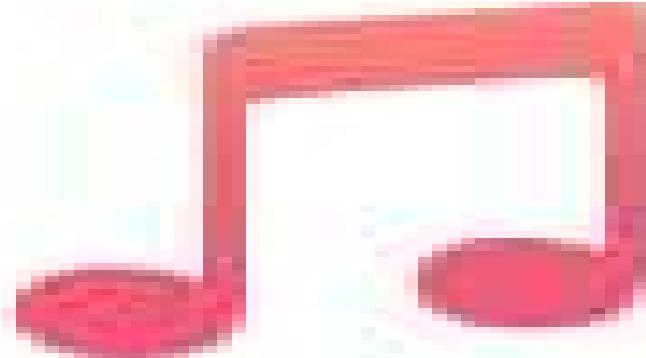
Switch: <https://arxiv.org/pdf/2101.03961.pdf>

Pre-training vs downstream quality



Alternative: FSDP

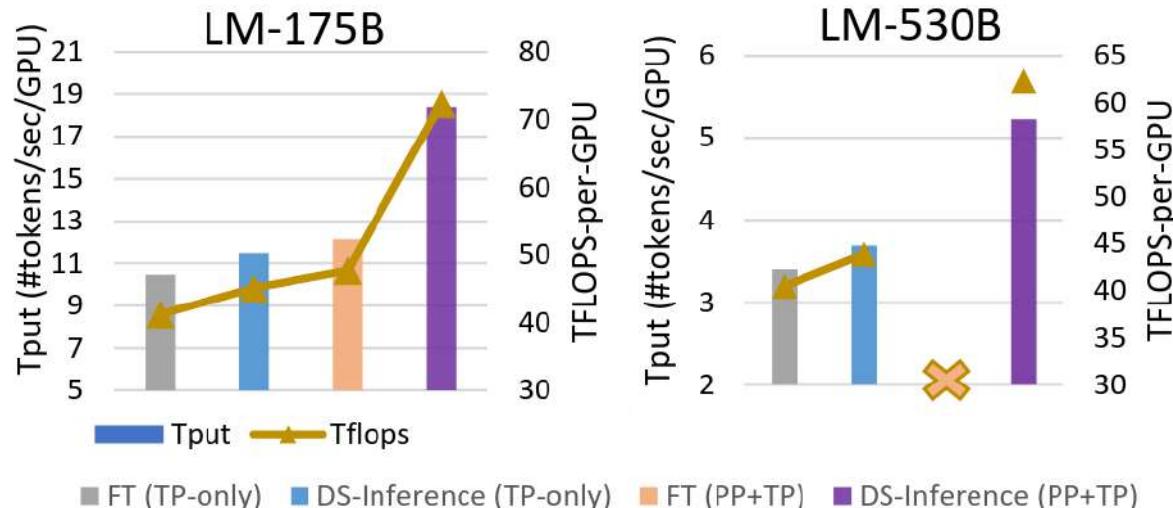
Source: [microsoft](#)



DeepSpeed Inference

Paper: <https://arxiv.org/abs/2207.00032>

- Same techniques, but for inference
- Offloading, tensor- & pipeline-parallel
- ... and a ton of hacks



</ZeRO>

Multi-GPU strategies:

- * Pipeline model-parallel – allocate layers on different GPUs
- * Sharded data-parallel – split optimizer state and/or parameters

Single GPU strategies:

- * Small model – gradient checkpointing & virtual batch
- * Large model – optimizer state sharding (keep parameters on GPU)

Implementations:

- DeepSpeed – sharded DP, offload, tensor parallelism, active development
 - Offload – <https://www.deepspeed.ai/news/2021/03/07/zero3-offload.html>
- FSDP – most of DeepSpeed features with native PyTorch API
- Model-specific implementations – <https://github.com/NVIDIA/Megatron-LM>

If we have time...
(if not, skip)

</lecture>

Example configuration:

Several GPU w/ 24GB memory | 128GB system memory | 16GBps interconnect

16GB model and optimizer, 128GB activations (batch 32) →

???

</lecture>

Example configuration:

Several GPU w/ 24GB memory | 128GB system memory | 16GBps interconnect

16GB model and optimizer, 128GB activations (batch 32) → **grad accumulation**

16GB model and optimizer, 16GB activations (batch 1) - ???

</lecture>

Example configuration:

Several GPU w/ 24GB memory | 128GB system memory | 16GBps interconnect

16GB model and optimizer, 128GB activations (batch 32) → **grad accumulation**

16GB model and optimizer, 16GB activations (batch 1) → **grad checkpointing**

32GB model and optimizer, 1GB activations → **???**

</lecture>

Example configuration:

Several GPU w/ 24GB memory | 128GB system memory | 16GBps interconnect

16GB model and optimizer, 128GB activations (batch 32) → **grad accumulation**

16GB model and optimizer, 16GB activations (batch 1) → **grad checkpointing**

32GB model and optimizer, 1GB activations → **it depends...**

DDP + offloading | **FSDP (ZeRO)** | **Pipeline-parallel** | **Tensor-parallel**
?

When is this the best option?

</lecture>

Example configuration:

Several GPU w/ 24GB memory | 128GB system memory | 16GBps interconnect

16GB model and optimizer, 128GB activations (batch 32) → **grad accumulation**

16GB model and optimizer, 16GB activations (batch 1) → **grad checkpointing**

32GB model and optimizer, 1GB activations → **it depends...**

DDP + offloading | **FSDP (ZeRO)** | Pipeline-parallel | Tensor-parallel

e.g. if too few GPUs
for other methods

?

When is this the best option?

</lecture>

Example configuration:

Several GPU w/ 24GB memory | 128GB system memory | 16GBps interconnect

16GB model and optimizer, 128GB activations (batch 32) → **grad accumulation**

16GB model and optimizer, 16GB activations (batch 1) → **grad checkpointing**

32GB model and optimizer, 1GB activations → **it depends...**

DDP + offloading | **FSDP (ZeRO)** | **Pipeline-parallel** | **Tensor-parallel**

*e.g. if too few GPUs
for other methods*

*no custom model code,
best for large batches*

?

When is this the best option?

</lecture>

Example configuration:

Several GPU w/ 24GB memory | 128GB system memory | 16GBps interconnect

16GB model and optimizer, 128GB activations (batch 32) → **grad accumulation**

16GB model and optimizer, 16GB activations (batch 1) → **grad checkpointing**

32GB model and optimizer, 1GB activations → **it depends...**

DDP + offloading | **FSDP (ZeRO)** | **Pipeline-parallel** | **Tensor-parallel**

e.g. if too few GPUs
for other methods

*no custom model code, communication-efficient
best for large batches sequential model*

?

When is this the best option?

</lecture>

Example configuration:

Several GPU w/ 24GB memory | 128GB system memory | 16GBps interconnect

16GB model and optimizer, 128GB activations (batch 32) → **grad accumulation**

16GB model and optimizer, 16GB activations (batch 1) → **grad checkpointing**

32GB model and optimizer, 1GB activations → **it depends...**

DDP + offloading | FSDP (ZeRO) | Pipeline-parallel | Tensor-parallel

e.g. if too few GPUs
for other methods

no custom model code, communication-efficient
best for large batches

minimal latency
non-symmetric model

Mix and match: TP within one server, minimal PP between servers, DDP between groups

Parallel code: manual (e.g. Megatron-LM) vs automated (alpa, FSDP, tensor_parallel)

Unconventional hardware: hivemind, petals, varuna, etc

</lecture>

Example configuration:

Several GPU w/ 24GB memory | 128GB system memory | 16GBps interconnect

16GB model and optimizer, 128GB activations (batch 32) → **grad accumulation**

16GB model and optimizer, 16GB activations (batch 1) → **grad checkpointing**

32GB model and optimizer, 1GB activations → **it depends...**

If the model does not fit, you can also quantize it into submission!
(more on model compression in a future lecture)



That's all Folks!